

# Tidy Time Series & Forecasting in R



## 6. Automatic forecasting algorithms

# Outline

- 1 What can we forecast?
- 2 Some case studies
- 3 The statistical forecasting perspective
- 4 Exponential smoothing
- 5 Lab Session 7
- 6 ARIMA models
- 7 Lab Session 8
- 8 Seasonal ARIMA models

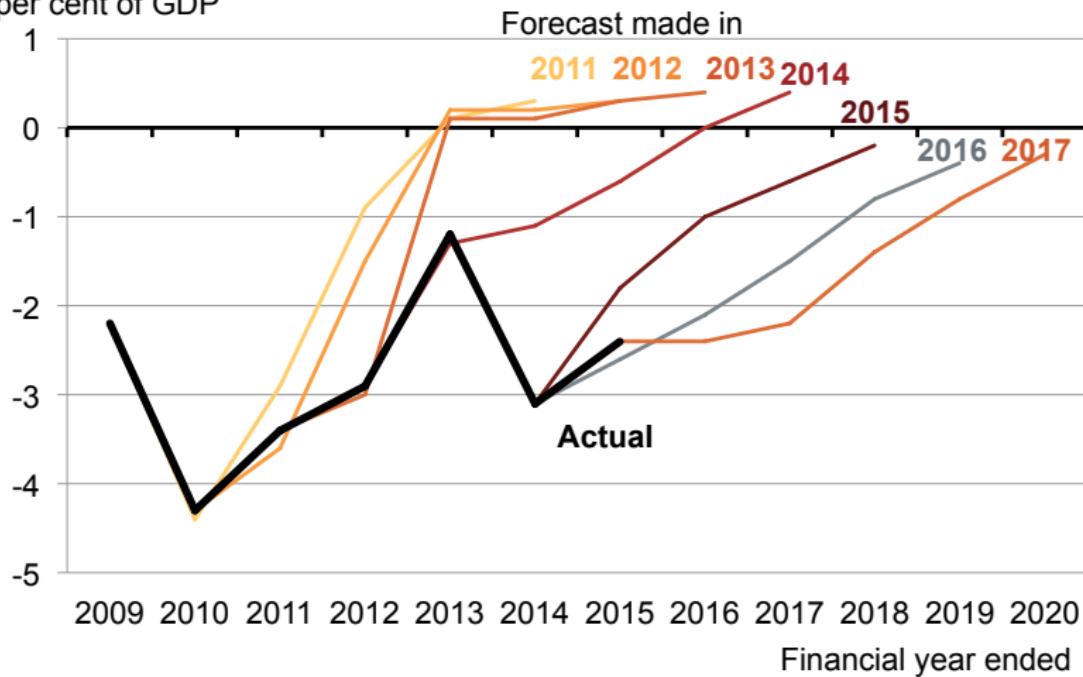
# Outline

- 1 What can we forecast?
- 2 Some case studies
- 3 The statistical forecasting perspective
- 4 Exponential smoothing
- 5 Lab Session 7
- 6 ARIMA models
- 7 Lab Session 8
- 8 Seasonal ARIMA models

# Forecasting is difficult

Commonwealth plans to drift back to surplus **GRATTAN Institute**  
show the triumph of experience over hope

Actual and forecast Commonwealth underlying cash balance  
per cent of GDP



# Forecasting is difficult

## A Timeline of Very Bad Future Predictions

1800



“ Rail travel at high speed is not possible, because passengers, unable to breathe, would die of asphyxia.”

Dr. Dionysus Larder, Professor of Natural Philosophy & Astronomy, University College London

1859



“ Drill for oil? You mean drill into the ground to try and find oil? You’re crazy!”

Associates of Edwin L. Drake refusing his suggestion to drill for oil in 1859 (Later that year, Drake succeeded in drilling the first oil well.)

1876



“ This telephone has too many shortcomings to be seriously considered as a means of communication.”

Western Union internal memo

1880



“ Everyone acquainted with the subject will recognize it as a conspicuous failure.”

Henry Morton, president of the Stevens Institute of Technology, on Edison's light bulb

1916



“ The idea that cavalry will be replaced by these iron coaches is absurd. It is little short of treasonous.”

Comment of Aide-de-camp to Field Marshal Haig, at tank demonstration

1946



“ Television won't last because people will soon get tired of staring at a plywood box every night.”

Darryl Zanuck, movie producer, 20th Century Fox

1902



“ Flight by machines heavier than air is unpractical and insignificant, if not utterly impossible.”

Simon Newcomb, Canadian-American astronomer and mathematician, 18 months before the Wright Brothers' flight at Kittyhawk

1916



“ The cinema is little more than a fad. It's canned drama. What audiences really want to see is flesh and blood on the stage.”

Charlie Chaplin, actor, producer, director, and studio founder

1977



“ There is no reason for any individual to have a computer in his home.”

Ken Olson, president, chairman and founder of Digital Equipment Corporation

1903



“ The horse is here to stay, but the automobile is only a novelty, a fad.”

The president of the Michigan Savings Bank, advising Henry Ford's lawyer not to invest in the Ford Motor Company

1921



“ The wireless music box has no imaginable commercial value. Who would pay for a message sent to no one in particular?”

Associates of commercial radio and television pioneer, David Sarnoff, responding to his call for investment in the radio



Read  
Newspapers  
Online



“ The truth is no online database will replace your daily newspaper.”

Clifford Stoll, Newsweek article entitled *The Internet? Bah!*

## What can we forecast?



# What can we forecast?



# What can we forecast?

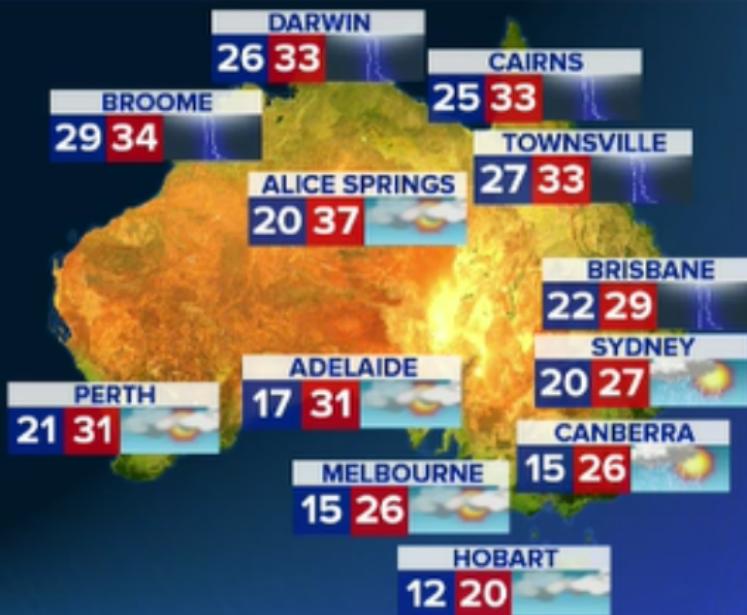


# What can we forecast?



# What can we forecast?

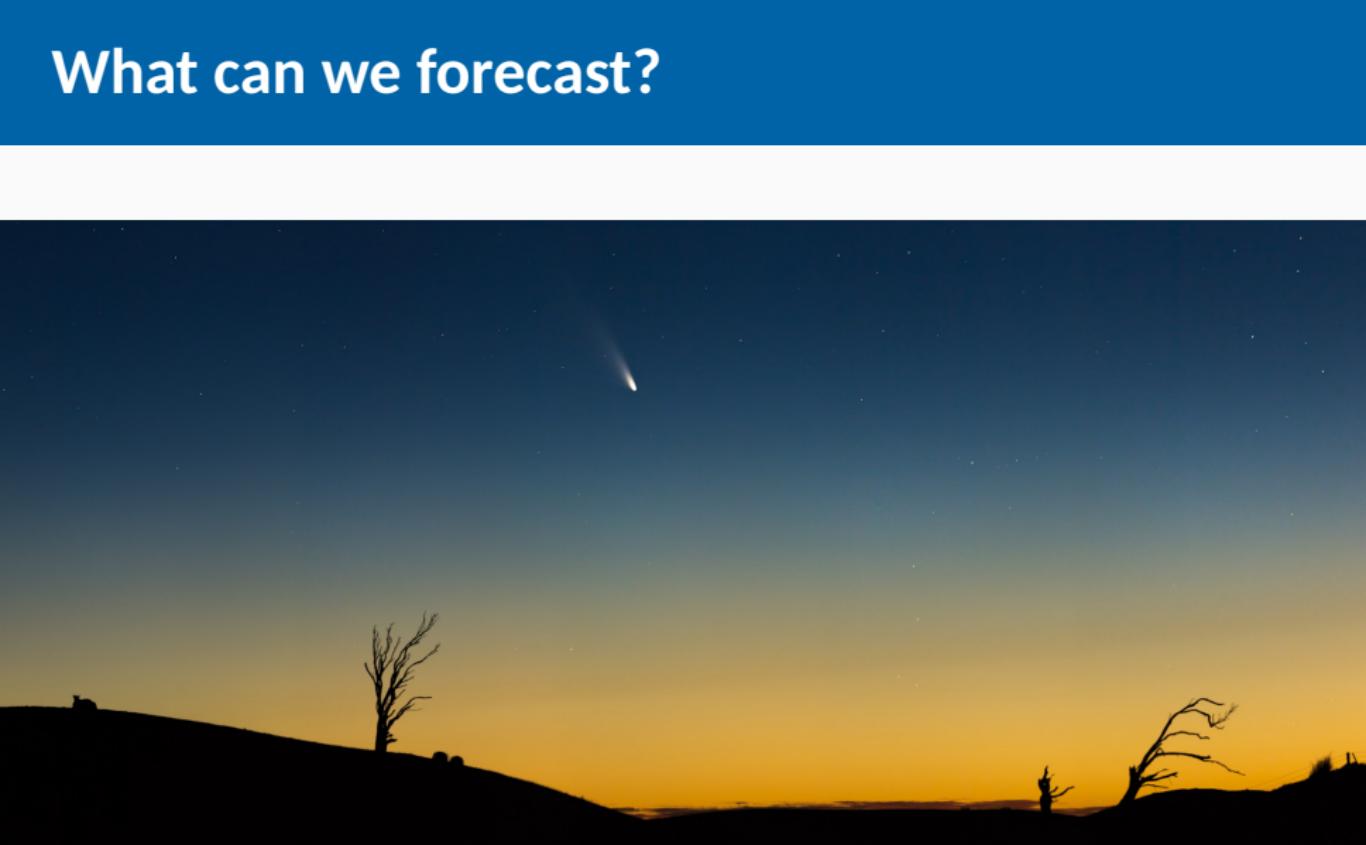
## TOMORROW



# What can we forecast?



# What can we forecast?



# Which is easiest to forecast?

- 1 daily electricity demand in 3 days time
- 2 timing of next Halley's comet appearance
- 3 time of sunrise this day next year
- 4 Google stock price tomorrow
- 5 Google stock price in 6 months time
- 6 maximum temperature tomorrow
- 7 exchange rate of \$US/AUS next week
- 8 total sales of drugs in Australian pharmacies next month

# Which is easiest to forecast?

- 1 daily electricity demand in 3 days time
  - 2 timing of next Halley's comet appearance
  - 3 time of sunrise this day next year
  - 4 Google stock price tomorrow
  - 5 Google stock price in 6 months time
  - 6 maximum temperature tomorrow
  - 7 exchange rate of \$US/AUS next week
  - 8 total sales of drugs in Australian pharmacies next month
- 
- how do we measure “easiest”?
  - what makes something easy/difficult to forecast?

# Factors affecting forecastability

Something is easier to forecast if:

- we have a good understanding of the factors that contribute to it
- there is lots of data available;
- the forecasts cannot affect the thing we are trying to forecast.
- there is relatively low natural/unexplainable random variation.
- the future is somewhat similar to the past

# Improving forecasts

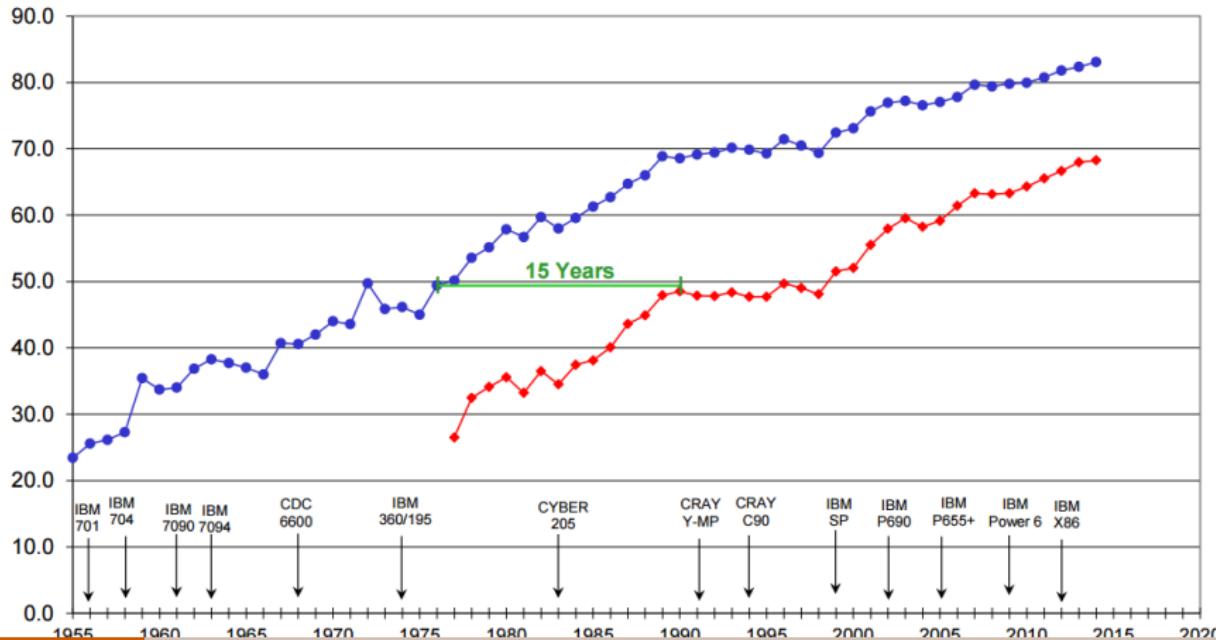


## NCEP Operational Forecast Skill

36 and 72 Hour Forecasts @ 500 MB over North America  
[ $100 * (1 - S1/70)$  Method]



—●— 36 Hour Forecast      —●— 72 Hour Forecast



# Outline

- 1 What can we forecast?
- 2 Some case studies
- 3 The statistical forecasting perspective
- 4 Exponential smoothing
- 5 Lab Session 7
- 6 ARIMA models
- 7 Lab Session 8
- 8 Seasonal ARIMA models

# CASE STUDY 1: Paperware company

**Problem:** Want forecasts of each of hundreds of items. Series can be stationary, trended or seasonal. They currently have a large forecasting program written in-house but it doesn't seem to produce sensible forecasts. They want me to tell them what is wrong and fix it.



## Additional information

- Program written in COBOL making numerical calculations limited. It is not possible to do any optimisation.
- Their programmer has little experience in numerical computing.
- They employ no statisticians and want

# CASE STUDY 1: Paperware company

## Methods currently used

- A 12 month average
- C 6 month average
- E straight line regression over last 12 months
- G straight line regression over last 6 months
- H average slope between last year's and this year's values. (Equivalent to differencing at lag 12 and taking mean.)
- I Same as H except over 6 months.
- K I couldn't understand the explanation.

## CASE STUDY 2: PBS



## CASE STUDY 2: PBS

**The Pharmaceutical Benefits Scheme (PBS) is the Australian government drugs subsidy scheme.**

- Many drugs bought from pharmacies are subsidised to allow more equitable access to modern drugs.
- The cost to government is determined by the number and types of drugs purchased. Currently nearly 1% of GDP.
- The total cost is budgeted based on forecasts of drug usage.

# CASE STUDY 2: PBS



# ABC News Online

AUSTRALIAN BROADCASTING CORPORATION



NewsRadio  
Streaming audio news  
LISTEN: [WMP](#) | [Real](#)

Select a Topic  
from the list below

- ▶ [Top Stories](#)
- ▶ [Just In](#)
- ▶ [World](#)
- ▶ [Asia-Pacific](#)
- ▶ [Business](#)
- ▶ [Sport](#)
- ▶ [Arts](#)
- ▶ [Sci Tech](#)
- ▶ [Indigenous](#)
- ▶ [Weather](#)
- ▶ [Rural](#)
- ▶ [Local News](#)
- ▶ [Broadband](#)

[Search](#)

## SPECIALS

- ▶ [Federal Election](#)

Click "Refresh" or "Reload"  
on your browser for the latest edition.

## POLITICS

### Opp demands drug price restriction after PBS budget blow-out

The Federal Opposition has called for tighter controls on drug prices after the Pharmaceutical Benefits Scheme (PBS) budget blew out by almost \$800 million.

The money was spent on two new drugs including the controversial anti-smoking aid Zyban, which dropped in price from \$220 to \$22 after it was listed on the PBS.

This Bulletin: Wed, May 30 2001 6:22 PM AEST



[For a fresh perspective on the federal election, reach into ABC Online's campaign weblog, The Poll Vault.](#)

[Audio News Online](#)

## CASE STUDY 2: PBS

- In 2001: \$4.5 billion budget, under-forecasted by \$800 million.
- Thousands of products. Seasonal demand.
- Subject to covert marketing, volatile products, uncontrollable expenditure.
- Although monthly data available for 10 years, data are aggregated to annual values, and only the first three years are used in estimating the forecasts.
- All forecasts being done with the FORECAST function in MS-Excel!

## CASE STUDY 3: Car fleet company

**Client:** One of Australia's largest car fleet companies

**Problem:** how to forecast resale value of vehicles?

How should this affect leasing and sales policies?

## CASE STUDY 3: Car fleet company

**Client:** One of Australia's largest car fleet companies

**Problem:** how to forecast resale value of vehicles?

How should this affect leasing and sales policies?

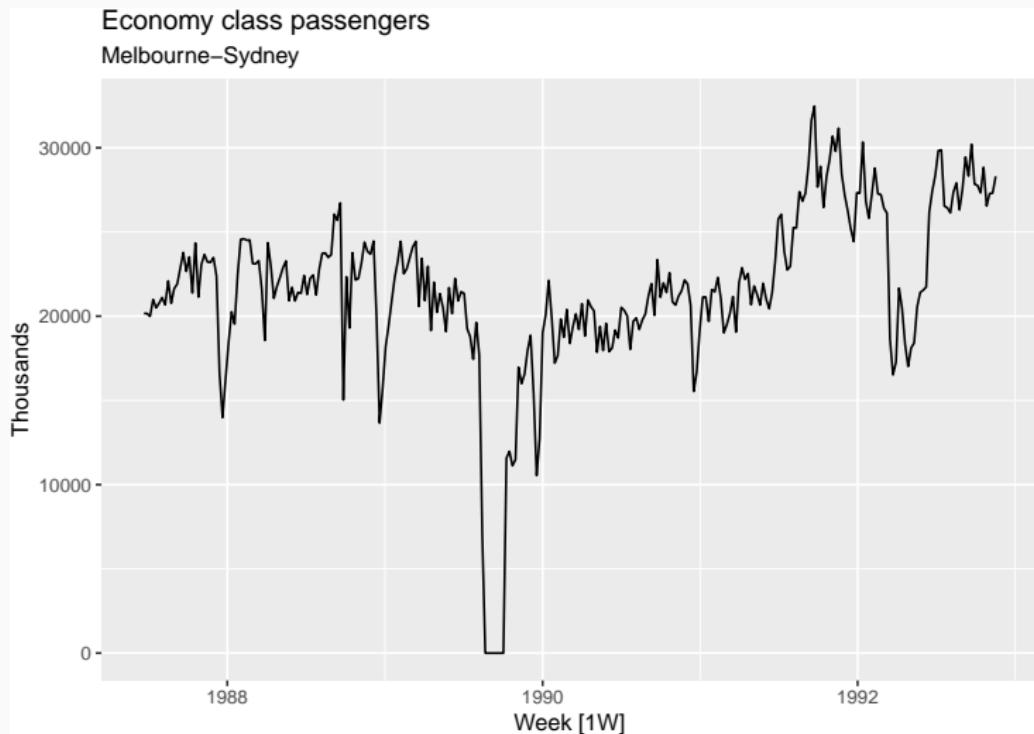
### Additional information

- They can provide a large amount of data on previous vehicles and their eventual resale values.
- The resale values are currently estimated by a group of specialists. They see me as a threat and do not cooperate.

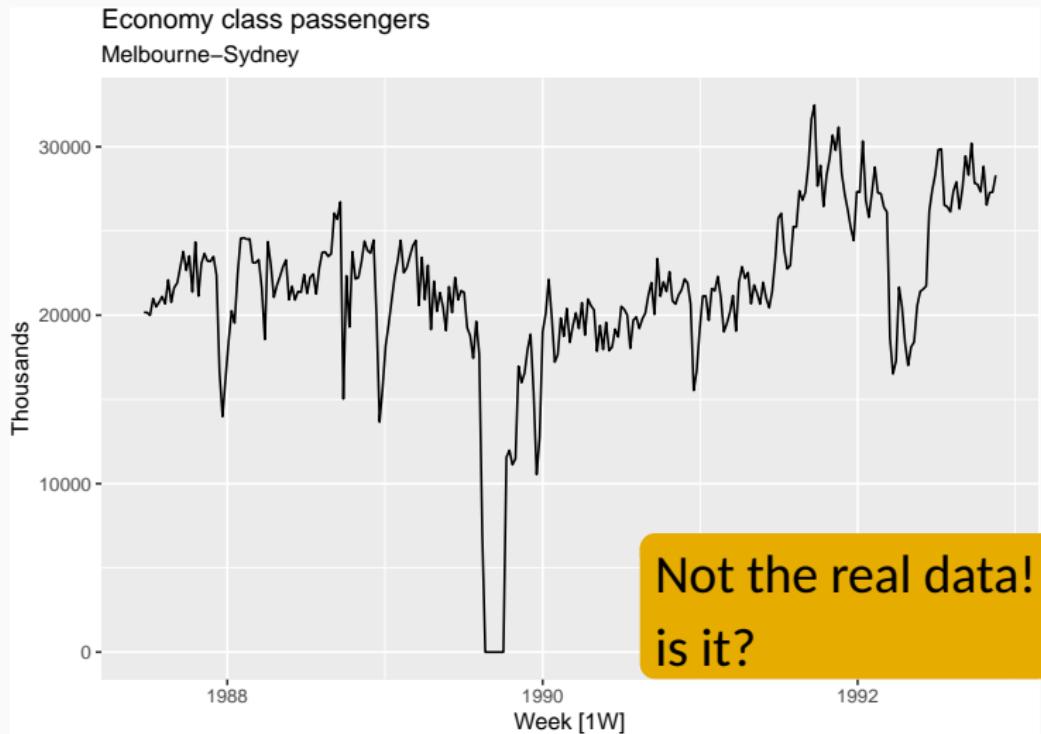
## CASE STUDY 4: Airline



# CASE STUDY 4: Airline



# CASE STUDY 4: Airline



# CASE STUDY 4: Airline

**Problem:** how to forecast passenger traffic on major routes?

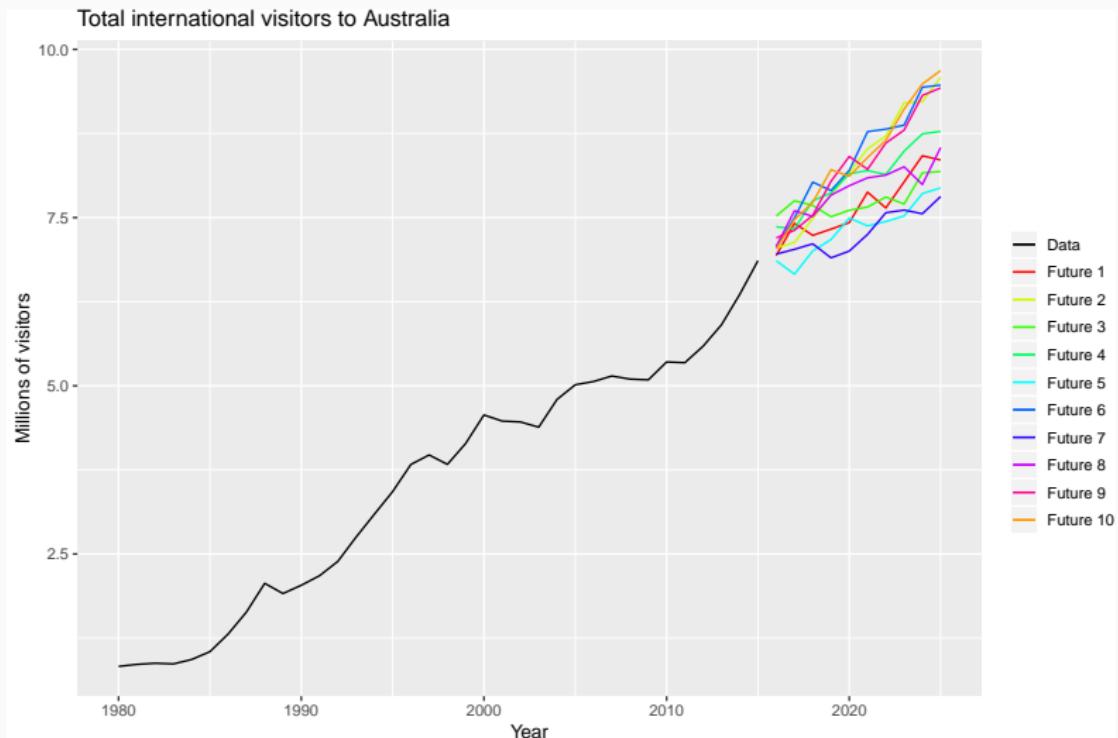
## Additional information

- They can provide a large amount of data on previous routes.
- Traffic is affected by school holidays, special events such as the Grand Prix, advertising campaigns, competition behaviour, etc.
- They have a highly capable team of people who are able to do most of the computing.

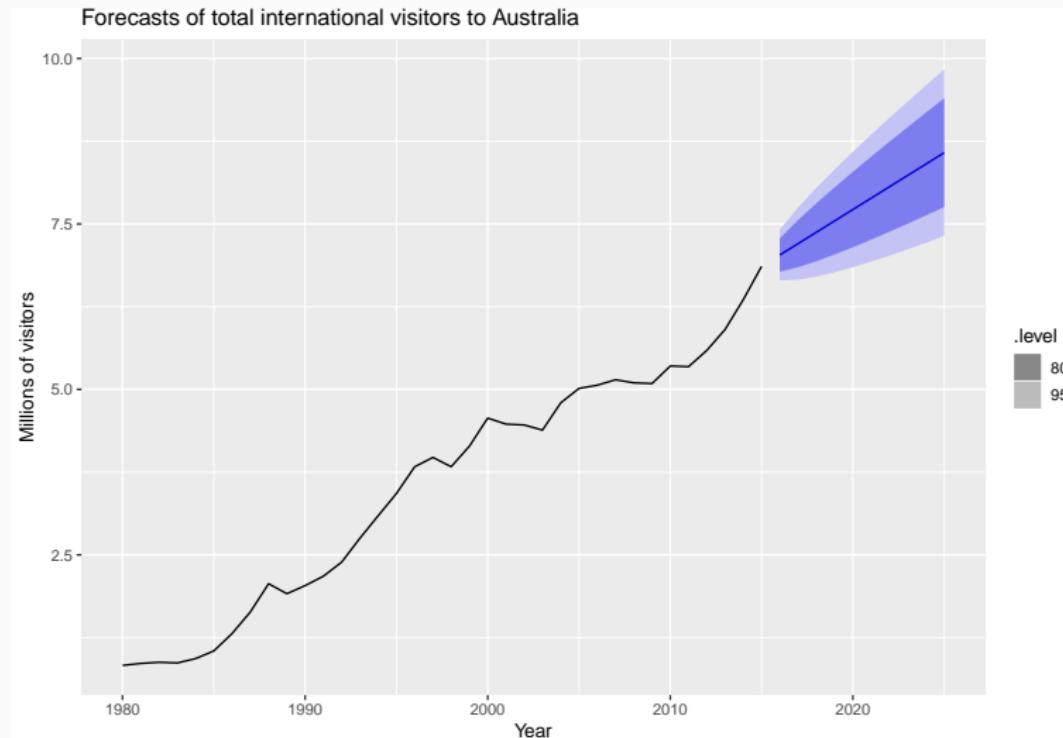
# Outline

- 1 What can we forecast?
- 2 Some case studies
- 3 The statistical forecasting perspective
- 4 Exponential smoothing
- 5 Lab Session 7
- 6 ARIMA models
- 7 Lab Session 8
- 8 Seasonal ARIMA models

# Sample futures



# Forecast intervals



# Statistical forecasting

- Thing to be forecast: a random variable,  $y_t$ .
- Forecast distribution: If  $\mathcal{I}$  is all observations, then  $y_t|\mathcal{I}$  means “the random variable  $y_t$  given what we know in  $\mathcal{I}$ ”.
- The “point forecast” is the mean (or median) of  $y_t|\mathcal{I}$
- The “forecast variance” is  $\text{var}[y_t|\mathcal{I}]$
- A prediction interval or “interval forecast” is a range of values of  $y_t$  with high probability.
- With time series,  $y_{t|t-1} = y_t|\{y_1, y_2, \dots, y_{t-1}\}$ .
- $\hat{y}_{T+h|T} = E[y_{T+h}|y_1, \dots, y_T]$  (an  $h$ -step forecast taking account of all observations up to time  $T$ ).

# Outline

- 1 What can we forecast?
- 2 Some case studies
- 3 The statistical forecasting perspective
- 4 Exponential smoothing
- 5 Lab Session 7
- 6 ARIMA models
- 7 Lab Session 8
- 8 Seasonal ARIMA models

# Historical perspective

- Developed in the 1950s and 1960s as methods (algorithms) to produce point forecasts.
- Combine a “level”, “trend” (slope) and “seasonal” component to describe a time series.
- The rate of change of the components are controlled by “smoothing parameters”:  $\alpha$ ,  $\beta$  and  $\gamma$  respectively.
- Need to choose best values for the smoothing parameters (and initial states).
- Equivalent ETS state space models developed in the 1990s and 2000s.

# A model for levels, trends, and seasonalities

We want a model that captures the level ( $\ell_t$ ), trend ( $b_t$ ) and seasonality ( $s_t$ ).

**How do we combine these elements?**

# A model for levels, trends, and seasonalities

We want a model that captures the level ( $\ell_t$ ), trend ( $b_t$ ) and seasonality ( $s_t$ ).

**How do we combine these elements?**

**Additively?**

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

# A model for levels, trends, and seasonalities

We want a model that captures the level ( $\ell_t$ ), trend ( $b_t$ ) and seasonality ( $s_t$ ).

**How do we combine these elements?**

**Additively?**

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

**Multiplicatively?**

$$y_t = \ell_{t-1} b_{t-1} s_{t-m} (1 + \varepsilon_t)$$

# A model for levels, trends, and seasonalities

We want a model that captures the level ( $\ell_t$ ), trend ( $b_t$ ) and seasonality ( $s_t$ ).

**How do we combine these elements?**

**Additively?**

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

**Multiplicatively?**

$$y_t = \ell_{t-1} b_{t-1} s_{t-m} (1 + \varepsilon_t)$$

**Perhaps a mix of both?**

$$y_t = (\ell_{t-1} + b_{t-1}) s_{t-m} + \varepsilon_t$$

# A model for levels, trends, and seasonalities

We want a model that captures the level ( $\ell_t$ ), trend ( $b_t$ ) and seasonality ( $s_t$ ).

**How do we combine these elements?**

**Additively?**

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

**Multiplicatively?**

$$y_t = \ell_{t-1} b_{t-1} s_{t-m} (1 + \varepsilon_t)$$

**Perhaps a mix of both?**

$$y_t = (\ell_{t-1} + b_{t-1}) s_{t-m} + \varepsilon_t$$

**How do the level, trend and seasonal components evolve over time?**

# ETS models

General notation

ETS : ExponenTial Smoothing  
Error Trend Season

**Error:** Additive ("A") or multiplicative ("M")

# ETS models

General notation

ETS : ExponenTial Smoothing  
Error Trend Season

**Error:** Additive ("A") or multiplicative ("M")

**Trend:** None ("N"), additive ("A"), multiplicative ("M"), or damped ("Ad" or "Md").

# ETS models

General notation

ETS : ExponenTial Smoothing  
Error Trend Season

**Error:** Additive ("A") or multiplicative ("M")

**Trend:** None ("N"), additive ("A"), multiplicative ("M"), or damped ("Ad" or "Md").

**Seasonality:** None ("N"), additive ("A") or multiplicative ("M")

## ETS(A,N,N): SES with additive errors

Forecast equation       $\hat{y}_{T+h|T} = \ell_T$

Measurement equation       $y_t = \ell_{t-1} + \varepsilon_t$

State equation       $\ell_t = \ell_{t-1} + \alpha \varepsilon_t$

where  $\varepsilon_t \sim \text{NID}(0, \sigma^2)$ .

## ETS(A,N,N): SES with additive errors

Forecast equation       $\hat{y}_{T+h|T} = \ell_T$

Measurement equation       $y_t = \ell_{t-1} + \varepsilon_t$

State equation       $\ell_t = \ell_{t-1} + \alpha \varepsilon_t$

where  $\varepsilon_t \sim \text{NID}(0, \sigma^2)$ .

- “innovations” or “single source of error” because equations have the same error process,  $\varepsilon_t$ .
- Measurement equation: relationship between observations and states.
- Transition/state equation(s): evolution of the state(s) through time.

## ETS(M,N,N): SES with multiplicative errors

Forecast equation       $\hat{y}_{T+h|T} = \ell_T$

Measurement equation       $y_t = \ell_{t-1}(1 + \varepsilon_t)$

State equation       $\ell_t = \ell_{t-1}(1 + \alpha\varepsilon_t)$

where  $\varepsilon_t \sim \text{NID}(0, \sigma^2)$ .

## ETS(M,N,N): SES with multiplicative errors

Forecast equation  $\hat{y}_{T+h|T} = \ell_T$

Measurement equation  $y_t = \ell_{t-1}(1 + \varepsilon_t)$

State equation  $\ell_t = \ell_{t-1}(1 + \alpha\varepsilon_t)$

where  $\varepsilon_t \sim \text{NID}(0, \sigma^2)$ .

- Models with additive and multiplicative errors with the same parameters generate the same point forecasts but different prediction intervals.

# ETS(A,A,N): Holt's linear trend

## Additive errors

Forecast equation  $\hat{y}_{T+h|T} = \ell_T + h b_T$

Measurement equation  $y_t = \ell_{t-1} + b_{t-1} + \varepsilon_t$

State equations  $\ell_t = \ell_{t-1} + b_{t-1} + \alpha \varepsilon_t$

$b_t = b_{t-1} + \beta \varepsilon_t$

# ETS(A,A,N): Holt's linear trend

## Additive errors

Forecast equation  $\hat{y}_{T+h|T} = \ell_T + hb_T$

Measurement equation  $y_t = \ell_{t-1} + b_{t-1} + \varepsilon_t$

State equations  $\ell_t = \ell_{t-1} + b_{t-1} + \alpha\varepsilon_t$

$b_t = b_{t-1} + \beta\varepsilon_t$

## Multiplicative errors

Forecast equation  $\hat{y}_{T+h|T} = \ell_T + hb_T$

Measurement equation  $y_t = (\ell_{t-1} + b_{t-1})(1 + \varepsilon_t)$

State equations  $\ell_t = (\ell_{t-1} + b_{t-1})(1 + \alpha\varepsilon_t)$

$b_t = b_{t-1} + \beta\varepsilon_t$

# Example: Australian population

```
aus_economy <- global_economy %>% filter(Code == "AUS") %>%
  mutate(Pop = Population/1e6)
fit <- aus_economy %>% model(AAN = ETS(Pop))
report(fit)
```

```
## Series: Pop
## Model: ETS(A,A,N)
##   Smoothing parameters:
##     alpha = 1
##     beta  = 0.327
##
##   Initial states:
##     l     b
## 10.1 0.222
##
##   sigma^2:  0.0041
##
##   AIC   AICc    BIC
## -77.0 -75.8 -66.7
```

# Example: Australian population

**components(fit)**

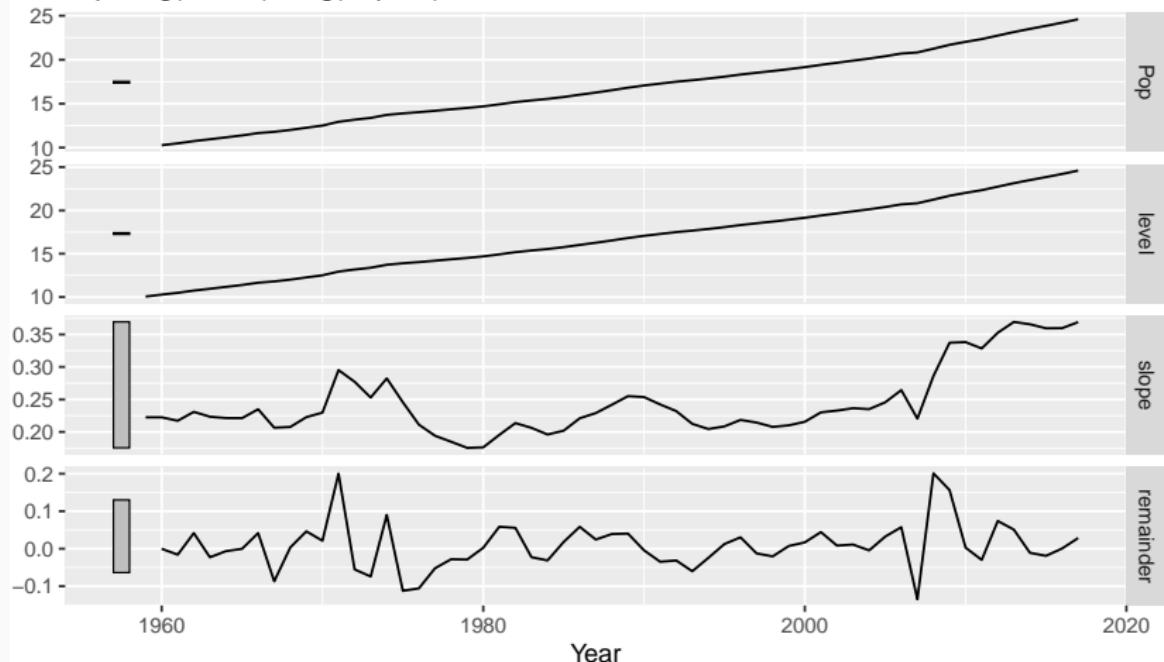
```
## # A dable:                      59 x 7 [1Y]
## # Key:                           Country, .model [1]
## # ETS(A,A,N) Decomposition: Pop = lag(level, 1) + lag(slope, 1)
## # remainder
##   Country   .model   Year   Pop level slope remainder
##   <fct>     <chr>    <dbl>  <dbl> <dbl>  <dbl>
## 1 Australia AAN     1959    NA    10.1  0.222  NA
## 2 Australia AAN     1960    10.3   10.3  0.222 -0.000145
## 3 Australia AAN     1961    10.5   10.5  0.217 -0.0159
## 4 Australia AAN     1962    10.7   10.7  0.231  0.0418
## 5 Australia AAN     1963    11.0   11.0  0.223 -0.0229
## 6 Australia AAN     1964    11.2   11.2  0.221 -0.00641
## 7 Australia AAN     1965    11.4   11.4  0.221 -0.000314
## 8 Australia AAN     1966    11.7   11.7  0.235  0.0418
## 9 Australia AAN     1967    11.8   11.8  0.206 -0.0869
## 10 Australia AAN    1968    12.0   12.0  0.209  0.00350
```

# Example: Australian population

```
components(fit) %>% autoplot()
```

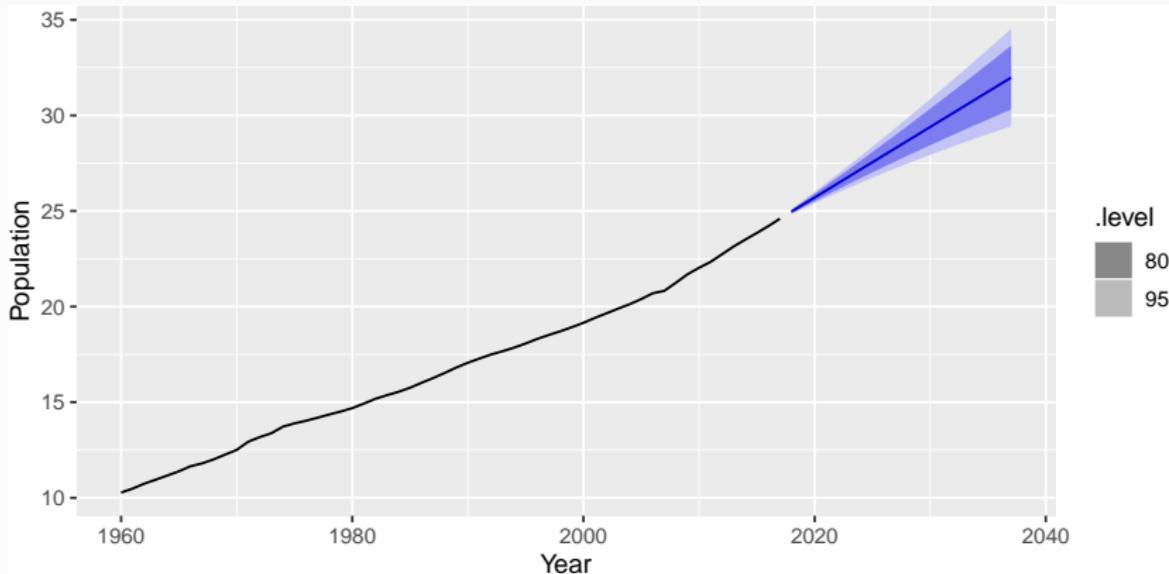
ETS(A,A,N) decomposition

Pop = lag(level, 1) + lag(slope, 1) + remainder



# Example: Australian population

```
fit %>%  
  forecast(h = 20) %>%  
  autoplot(aus_economy) +  
  ylab("Population") + xlab("Year")
```



# ETS(A,Ad,N): Damped trend method

## Additive errors

Forecast equation  $\hat{y}_{T+h|T} = \ell_T + (\phi + \dots + \phi^{h-1})b_T$

Measurement equation  $y_t = (\ell_{t-1} + \phi b_{t-1}) + \varepsilon_t$

State equations  $\ell_t = (\ell_{t-1} + \phi b_{t-1}) + \alpha \varepsilon_t$

$b_t = \phi b_{t-1} + \beta \varepsilon_t$

# ETS(A,Ad,N): Damped trend method

## Additive errors

Forecast equation  $\hat{y}_{T+h|T} = \ell_T + (\phi + \dots + \phi^{h-1})b_T$

Measurement equation  $y_t = (\ell_{t-1} + \phi b_{t-1}) + \varepsilon_t$

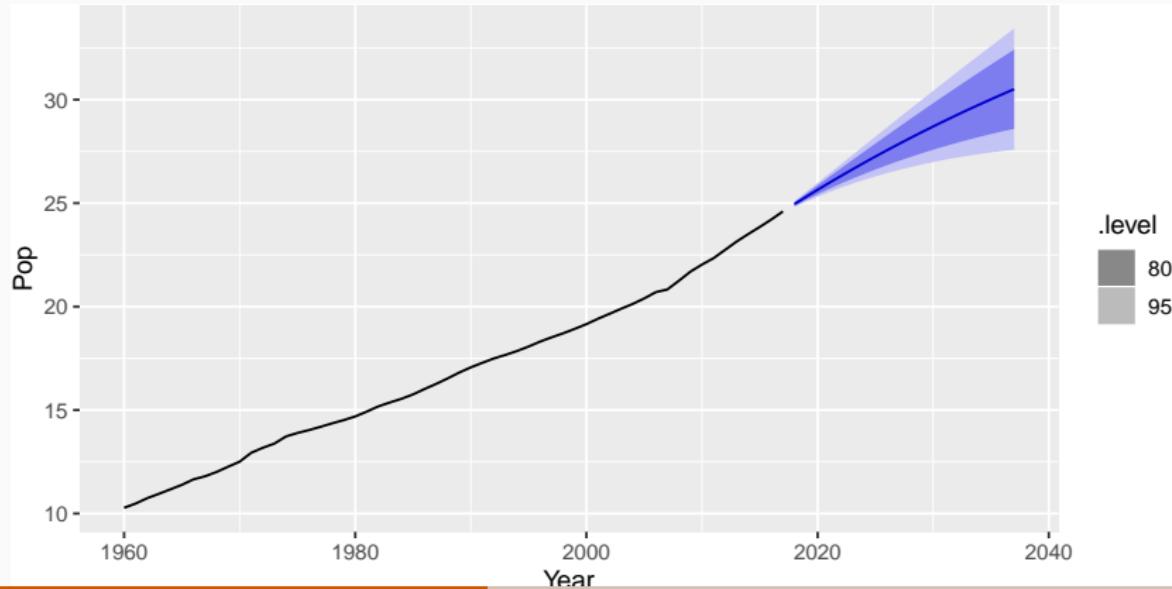
State equations  $\ell_t = (\ell_{t-1} + \phi b_{t-1}) + \alpha \varepsilon_t$

$b_t = \phi b_{t-1} + \beta \varepsilon_t$

- Damping parameter  $0 < \phi < 1$ .
- If  $\phi = 1$ , identical to Holt's linear trend.
- As  $h \rightarrow \infty$ ,  $\hat{y}_{T+h|T} \rightarrow \ell_T + \phi b_T / (1 - \phi)$ .
- Short-run forecasts trended, long-run forecasts constant.

# Example: Australian population

```
aus_economy %>%  
  model(holt = ETS(Pop ~ trend("Ad")))) %>%  
  forecast(h = 20) %>%  
  autoplot(aus_economy)
```



# Example: National populations

```
fit <- global_economy %>%  
  mutate(Pop = Population/1e6) %>%  
  model(ets = ETS(Pop))  
fit  
  
## # A mable: 263 x 2  
## # Key:      Country [263]  
##   Country           ets  
##   <fct>            <model>  
## 1 Afghanistan       <ETS(A,A,N)>  
## 2 Albania           <ETS(M,A,N)>  
## 3 Algeria           <ETS(M,A,N)>  
## 4 American Samoa    <ETS(M,A,N)>  
## 5 Andorra           <ETS(M,A,N)>  
## 6 Angola             <ETS(M,A,N)>  
## 7 Antigua and Barbuda <ETS(M,A,N)>  
## 8 Arab World          <ETS(M,A,N)>  
## 9 Argentina          <ETS(A,A,N)>  
## 10 Armenia            <ETS(M,A,N)>  
## # ... with 253 more rows
```

# Example: National populations

```
fit %>%  
  forecast(h = 5)
```

```
## # A fable: 1,315 x 5 [1Y]  
## # Key:      Country, .model [263]  
##   Country     .model Year   Pop .distribution  
##   <fct>       <chr>  <dbl> <dbl> <dist>  
## 1 Afghanistan ets    2018 36.4 N(36, 0.012)  
## 2 Afghanistan ets    2019 37.3 N(37, 0.059)  
## 3 Afghanistan ets    2020 38.2 N(38, 0.164)  
## 4 Afghanistan ets    2021 39.0 N(39, 0.351)  
## 5 Afghanistan ets    2022 39.9 N(40, 0.644)  
## 6 Albania      ets    2018  2.87 N(2.9, 0.00012)  
## 7 Albania      ets    2019  2.87 N(2.9, 0.00060)  
## 8 Albania      ets    2020  2.87 N(2.9, 0.00169)  
## 9 Albania      ets    2021  2.86 N(2.9, 0.00362)
```

## ETS(A,A,A): Holt-Winters additive method

Forecast equation  $\hat{y}_{t+h|t} = \ell_t + hb_t + s_{t+h-m(k+1)}$

Observation equation  $y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$

State equations  $\ell_t = \ell_{t-1} + b_{t-1} + \alpha \varepsilon_t$

$b_t = b_{t-1} + \beta \varepsilon_t$

$s_t = s_{t-m} + \gamma \varepsilon_t$

- $k = \text{integer part of } (h - 1)/m.$
- $\sum_i s_i \approx 0.$
- Parameters:  $0 \leq \alpha \leq 1$ ,  $0 \leq \beta^* \leq 1$ ,  $0 \leq \gamma \leq 1 - \alpha$  and  $m = \text{period of seasonality}$  (e.g.  $m = 4$  for quarterly data).

## ETS(M,A,M): Holt-Winters multiplicative method

Forecast equation  $\hat{y}_{t+h|t} = (\ell_t + hb_t)s_{t+h-m(k+1)}$

Observation equation  $y_t = (\ell_{t-1} + b_{t-1})s_{t-m}(1 + \varepsilon_t)$

State equations  $\ell_t = (\ell_{t-1} + b_{t-1})(1 + \alpha\varepsilon_t)$

$$b_t = b_{t-1}(1 + \beta\varepsilon_t)$$

$$s_t = s_{t-m}(1 + \gamma\varepsilon_t)$$

- $k$  is integer part of  $(h - 1)/m$ .
- $\sum_i s_i \approx m$ .
- Parameters:  $0 \leq \alpha \leq 1$ ,  $0 \leq \beta^* \leq 1$ ,  $0 \leq \gamma \leq 1 - \alpha$  and  $m$  = period of seasonality (e.g.  $m = 4$  for quarterly data).

# Example: Australian holiday tourism

```
holidays <- tourism %>%
  filter(Purpose == "Holiday")
fit <- holidays %>% model(ets = ETS(Trips))
fit

## # A mable: 76 x 4
## # Key:      Region, State, Purpose [76]
##   Region                  State          Purpose ets
##   <chr>                   <chr>        <chr>   <model>
## 1 Adelaide                South Australia Holiday <ETS(A,N,A~
## 2 Adelaide Hills           South Australia Holiday <ETS(A,A,N~
## 3 Alice Springs            Northern Territory Holiday <ETS(M,N,A~
## 4 Australia's Coral Coast Western Australia Holiday <ETS(M,N,A~
## 5 Australia's Golden Outba~ Western Australia Holiday <ETS(M,N,M~
## 6 Australia's North West   Western Australia Holiday <ETS(A,N,A~
## 7 Australia's South West   Western Australia Holiday <ETS(M,N,M~
## 8 Ballarat                 Victoria       Holiday <ETS(M,N,A~
## 9 Barkly                  Northern Territory Holiday <ETS(A,N,A~
## 10 Barossa                 South Australia Holiday <ETS(A,N,N~
```

# Example: Australian holiday tourism

```
fit %>% filter(Region=="Snowy Mountains") %>% report()
```

```
## Series: Trips
## Model: ETS(M,N,A)
##   Smoothing parameters:
##     alpha = 0.157
##     gamma = 1e-04
##
##   Initial states:
##     l    s1    s2    s3    s4
##   142 -61  131 -42.2 -27.7
##
##   sigma^2:  0.0388
##
##   AIC  AICc   BIC
##   852   854   869
```

# Example: Australian holiday tourism

```
fit %>% filter(Region=="Snowy Mountains") %>% components(fit)
```

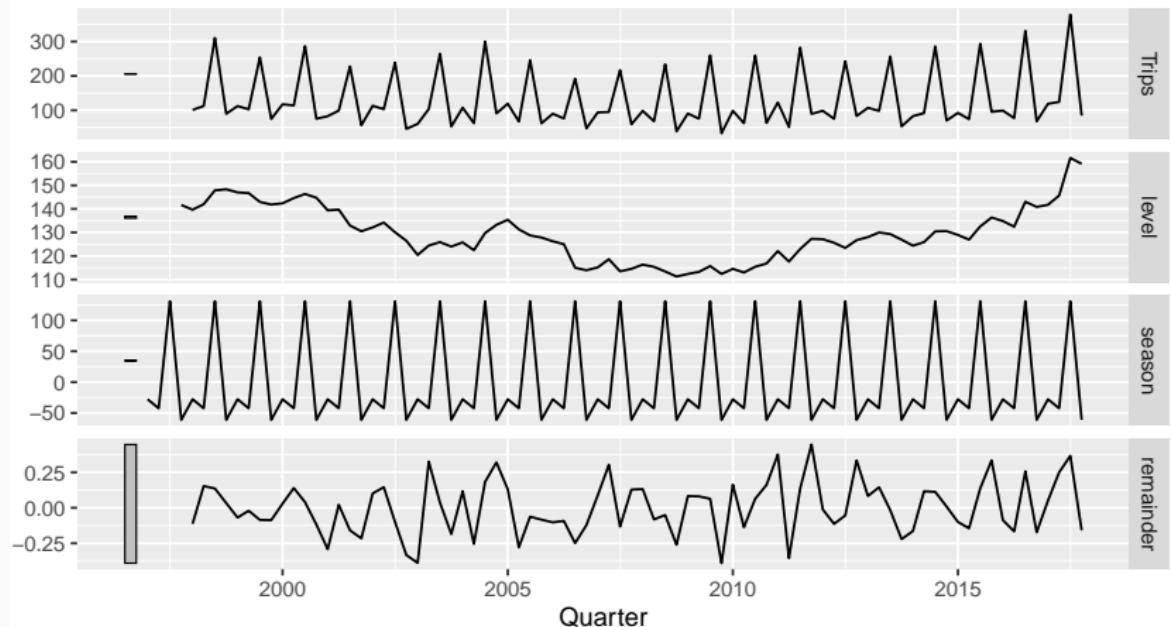
```
## # A dable:          84 x 9 [1Q]
## # Key:             Region, State, Purpose, .model [1]
## # ETS(M,N,A) Decomposition: Trips = (lag(level, 1) + lag(season,
## #     4)) * (1 + remainder)
## #   Region State Purpose .model      Quarter Trips level season
## #   <chr>   <chr>  <chr>    <chr>      <qtr>   <dbl> <dbl>   <dbl>
## # 1 Snowy~ New ~ Holiday ets      1997 Q1    NA     NA    -27.7
## # 2 Snowy~ New ~ Holiday ets      1997 Q2    NA     NA    -42.2
## # 3 Snowy~ New ~ Holiday ets      1997 Q3    NA     NA    131.
## # 4 Snowy~ New ~ Holiday ets      1997 Q4    NA    142.   -61.0
## # 5 Snowy~ New ~ Holiday ets      1998 Q1  101.   140.   -27.7
## # 6 Snowy~ New ~ Holiday ets      1998 Q2  112.   142.   -42.2
## # 7 Snowy~ New ~ Holiday ets      1998 Q3  310.   148.   131.
## # 8 Snowy~ New ~ Holiday ets      1998 Q4  89.8   148.   -61.0
## # 9 Snowy~ New ~ Holiday ets      1999 Q1  112.   147.   -27.7
## # 10 Snowy~ New ~ Holiday ets     1999 Q2  103.   147.   -42.2
## # ... with 74 more rows, and 1 more variable: remainder <dbl>
```

# Example: Australian holiday tourism

```
fit %>% filter(Region=="Snowy Mountains") %>%  
  components(fit) %>% autoplot()
```

ETS(M,N,A) decomposition

Trips =  $(\text{lag}(\text{level}, 1) + \text{lag}(\text{season}, 4)) * (1 + \text{remainder})$



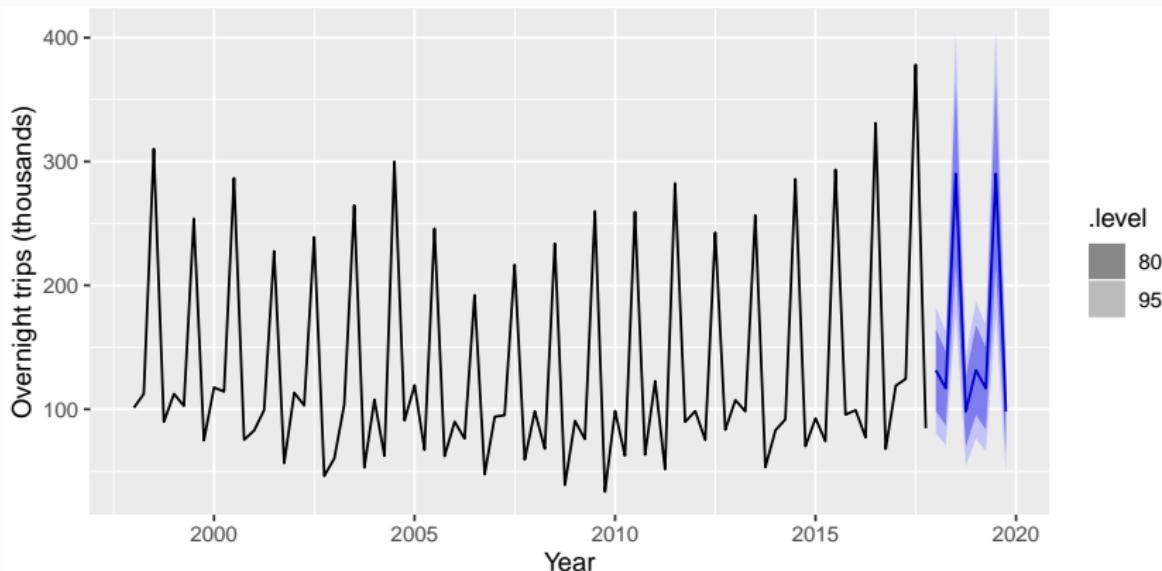
# Example: Australian holiday tourism

```
fit %>% forecast()

## # A fable: 608 x 7 [1Q]
## # Key:      Region, State, Purpose, .model [76]
##   Region     State    Purpose .model   Quarter Trips .distribution
##   <chr>      <chr>    <chr>   <chr>     <qtr> <dbl> <dist>
## 1 Adelaide  South A~ Holiday ets      2018 Q1 210. N(210, 457)
## 2 Adelaide  South A~ Holiday ets      2018 Q2 173. N(173, 473)
## 3 Adelaide  South A~ Holiday ets      2018 Q3 169. N(169, 489)
## 4 Adelaide  South A~ Holiday ets      2018 Q4 186. N(186, 505)
## 5 Adelaide  South A~ Holiday ets      2019 Q1 210. N(210, 521)
## 6 Adelaide  South A~ Holiday ets      2019 Q2 173. N(173, 537)
## 7 Adelaide  South A~ Holiday ets      2019 Q3 169. N(169, 553)
## 8 Adelaide  South A~ Holiday ets      2019 Q4 186. N(186, 569)
## 9 Adelaide~ South A~ Holiday ets      2018 Q1 19.4 N(19, 36)
## 10 Adelaide~ South A~ Holiday ets     2018 Q2 19.6 N(20, 36)
## # ... with 598 more rows
```

# Example: Australian holiday tourism

```
fit %>% forecast() %>%
  filter(Region=="Snowy Mountains") %>%
  autoplot(holidays) +
  xlab("Year") + ylab("Overnight trips (thousands)")
```



# Exponential smoothing models

Additive Error		Seasonal Component		
Trend Component		N (None)	A (Additive)	M (Multiplicative)
N	(None)	A,N,N	A,N,A	A,N,M
A	(Additive)	A,A,N	A,A,A	A,A,M
$A_d$	(Additive damped)	A, $A_d$ ,N	A, $A_d$ ,A	A, $A_d$ ,M

Multiplicative Error		Seasonal Component		
Trend Component		N (None)	A (Additive)	M (Multiplicative)
N	(None)	M,N,N	M,N,A	M,N,M
A	(Additive)	M,A,N	M,A,A	M,A,M
$A_d$	(Additive damped)	M, $A_d$ ,N	M, $A_d$ ,A	M, $A_d$ ,M

# Estimating ETS models

- Smoothing parameters  $\alpha, \beta, \gamma$  and  $\phi$ , and the initial states  $\ell_0, b_0, s_0, s_{-1}, \dots, s_{-m+1}$  are estimated by maximising the “likelihood” = the probability of the data arising from the specified model.
- For models with additive errors equivalent to minimising SSE.
- For models with multiplicative errors, **not** equivalent to minimising SSE.

# Model selection

## Akaike's Information Criterion

$$AIC = -2 \log(L) + 2k$$

where  $L$  is the likelihood and  $k$  is the number of parameters initial states estimated in the model.

# Model selection

## Akaike's Information Criterion

$$AIC = -2 \log(L) + 2k$$

where  $L$  is the likelihood and  $k$  is the number of parameters initial states estimated in the model.

## Corrected AIC

$$AIC_c = AIC + \frac{2(k+1)(k+2)}{T-k}$$

which is the AIC corrected (for small sample bias).

# Model selection

## Akaike's Information Criterion

$$AIC = -2 \log(L) + 2k$$

where  $L$  is the likelihood and  $k$  is the number of parameters initial states estimated in the model.

## Corrected AIC

$$AIC_c = AIC + \frac{2(k+1)(k+2)}{T-k}$$

which is the AIC corrected (for small sample bias).

## Bayesian Information Criterion

$$BIC = AIC + k(\log(T) - 2).$$

## AIC and cross-validation

Minimizing the AIC assuming Gaussian residuals is asymptotically equivalent to minimizing one-step time series cross validation MSE.

# Automatic forecasting

From Hyndman et al. (IJF, 2002):

- 1 Apply each model that is appropriate to the data.  
Optimize parameters and initial values using  
MLE.
  - 2 Select best method using AICc.
  - 3 Produce forecasts using best method.
  - 4 Obtain forecast intervals using underlying state  
space model.
- Method performed very well in M3 competition.
  - Used as a benchmark in the M4 competition.

# Outline

- 1 What can we forecast?
- 2 Some case studies
- 3 The statistical forecasting perspective
- 4 Exponential smoothing
- 5 Lab Session 7
- 6 ARIMA models
- 7 Lab Session 8
- 8 Seasonal ARIMA models

# Lab Session 7

Find an ETS model for the Gas data from aus\_production.

- Why is multiplicative seasonality necessary here?
- Experiment with making the trend damped.

# Outline

- 1 What can we forecast?
- 2 Some case studies
- 3 The statistical forecasting perspective
- 4 Exponential smoothing
- 5 Lab Session 7
- 6 ARIMA models
- 7 Lab Session 8
- 8 Seasonal ARIMA models

# ARIMA models

- AR:** autoregressive (lagged observations as inputs)
- I:** integrated (differencing to make series stationary)
- MA:** moving average (lagged errors as inputs)

# ARIMA models

- AR: autoregressive (lagged observations as inputs)
- I: integrated (differencing to make series stationary)
- MA: moving average (lagged errors as inputs)

An ARIMA model is rarely interpretable in terms of visible data structures like trend and seasonality. But it can capture a huge range of time series patterns.

# Stationarity

## Definition

If  $\{y_t\}$  is a stationary time series, then for all  $s$ , the distribution of  $(y_t, \dots, y_{t+s})$  does not depend on  $t$ .

# Stationarity

## Definition

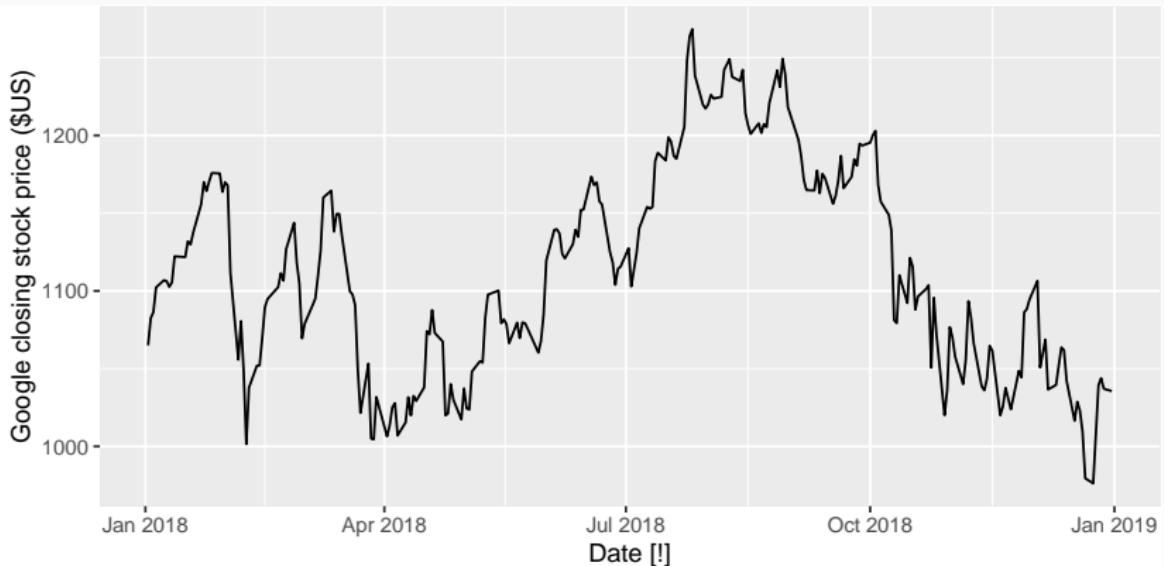
If  $\{y_t\}$  is a stationary time series, then for all  $s$ , the distribution of  $(y_t, \dots, y_{t+s})$  does not depend on  $t$ .

A stationary series is:

- roughly horizontal
- constant variance
- no patterns predictable in the long-term

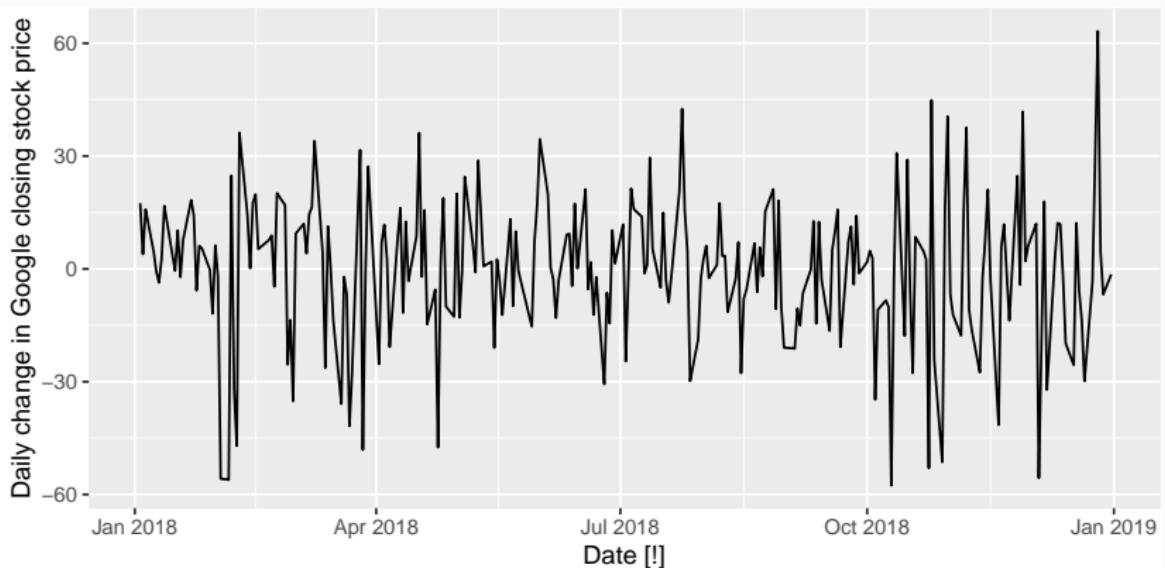
# Stationary?

```
gafa_stock %>%  
  filter(Symbol == "GOOG", year(Date) == 2018) %>%  
  autoplot(Close) +  
  ylab("Google closing stock price ($US)")
```



# Stationary?

```
gafa_stock %>%
  filter(Symbol == "GOOG", year(Date) == 2018) %>%
  autoplot(difference(Close)) +
  ylab("Daily change in Google closing stock price")
```



# Differencing

- Differencing helps to **stabilize the mean**.
- The differenced series is the *change* between each observation in the original series.
- Occasionally the differenced data will not appear stationary and it may be necessary to difference the data a second time.
- In practice, it is almost never necessary to go beyond second-order differences.

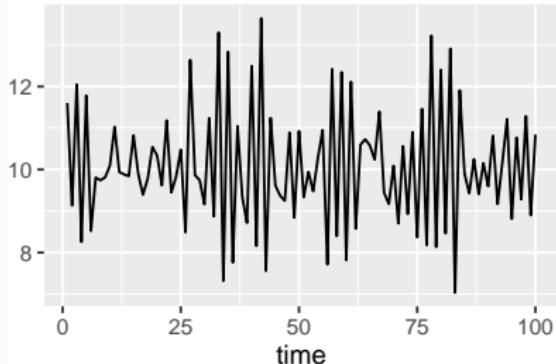
# Autoregressive models

## Autoregressive (AR) models:

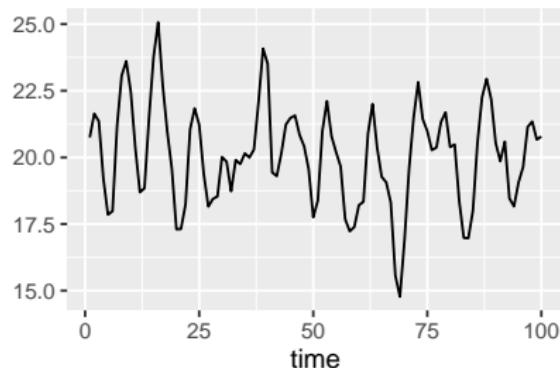
$$y_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \cdots + \phi_p y_{t-p} + \varepsilon_t,$$

where  $\varepsilon_t$  is white noise. This is a multiple regression with **lagged values** of  $y_t$  as predictors.

AR(1)



AR(2)



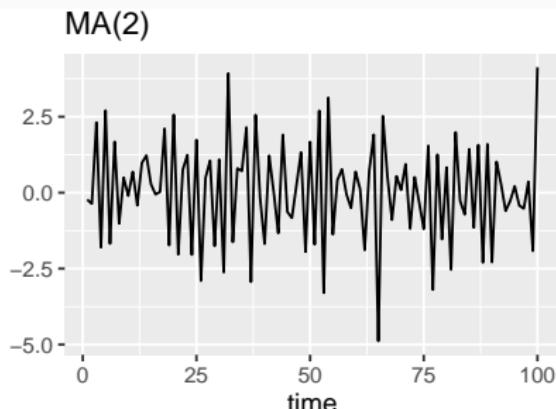
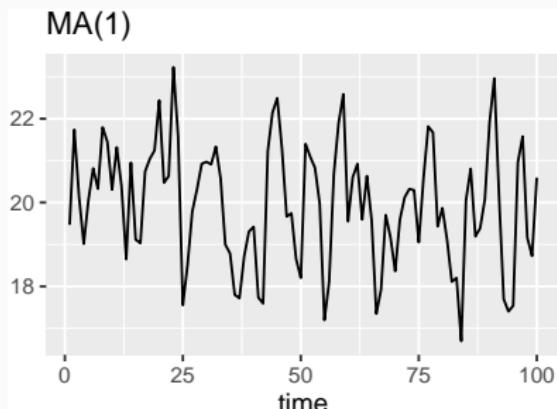
- Cyclic behaviour is possible when  $p \geq 2$ .

# Moving Average (MA) models

## Moving Average (MA) models:

$$y_t = c + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \cdots + \theta_q \varepsilon_{t-q},$$

where  $\varepsilon_t$  is white noise. This is a multiple regression with **lagged errors** as predictors. *Don't confuse this with moving average smoothing!*



# ARIMA models

## Autoregressive Moving Average models:

$$y_t = c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} \\ + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t.$$

# ARIMA models

## Autoregressive Moving Average models:

$$y_t = c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} \\ + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t.$$

- Predictors include both **lagged values of  $y_t$  and lagged errors.**

# ARIMA models

## Autoregressive Moving Average models:

$$y_t = c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} \\ + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t.$$

- Predictors include both **lagged values of  $y_t$  and lagged errors.**

## Autoregressive Integrated Moving Average models

- Combine ARMA model with **differencing**.
- $d$ -differenced series follows an ARMA model.
- Need to choose  $p, d, q$  and whether or not to include  $c$ .

# ARIMA models

## ARIMA( $p, d, q$ ) model

AR:  $p$  = order of the autoregressive part

I:  $d$  = degree of first differencing involved

MA:  $q$  = order of the moving average part.

- White noise model: ARIMA(0,0,0)
- Random walk: ARIMA(0,1,0) with no constant
- Random walk with drift: ARIMA(0,1,0) with const.
- AR( $p$ ): ARIMA( $p, 0, 0$ )
- MA( $q$ ): ARIMA( $0, 0, q$ )

# Example: National populations

```
fit <- global_economy %>%  
  model(arima = ARIMA(Population))  
fit  
  
## # A mable: 263 x 2  
## # Key:      Country [263]  
##   Country           arima  
##   <fct>             <model>  
## 1 Afghanistan       <ARIMA(4,2,1)>  
## 2 Albania            <ARIMA(0,2,2)>  
## 3 Algeria             <ARIMA(2,2,2)>  
## 4 American Samoa     <ARIMA(2,2,2)>  
## 5 Andorra            <ARIMA(2,1,2) w/ drift>  
## 6 Angola              <ARIMA(4,2,1)>  
## 7 Antigua and Barbuda <ARIMA(2,1,2) w/ drift>  
## 8 Arab World           <ARIMA(0,2,1)>
```

# Example: National populations

```
fit %>% filter(Country=="Australia") %>% report()

## Series: Population
## Model: ARIMA(0,2,1)
##
## Coefficients:
##         ma1
##         -0.661
## s.e.    0.107
##
## sigma^2 estimated as 4.063e+09: log likelihood=-699
## AIC=1401    AICc=1402    BIC=1405
```

# Example: National populations

```
fit %>% filter(Country=="Australia") %>% report()
```

```
## Series: Population
## Model: ARIMA(0,2,1)
##
## Coefficients:
##             ma1
##             -0.661
## s.e.      0.107
##
## sigma^2 estimated as 4.063e+09:  log likelihood=-699
## AIC=1401    AICc=1402    BIC=1405
```

$$y_t = 2y_{t-1} - y_{t-2} - 0.7\varepsilon_{t-1} + \varepsilon_t$$
$$\varepsilon_t \sim \text{NID}(0, 4 \times 10^9)$$

# Understanding ARIMA models

- If  $c = 0$  and  $d = 0$ , the long-term forecasts will go to zero.
- If  $c = 0$  and  $d = 1$ , the long-term forecasts will go to a non-zero constant.
- If  $c = 0$  and  $d = 2$ , the long-term forecasts will follow a straight line.
- If  $c \neq 0$  and  $d = 0$ , the long-term forecasts will go to the mean of the data.
- If  $c \neq 0$  and  $d = 1$ , the long-term forecasts will follow a straight line.
- If  $c \neq 0$  and  $d = 2$ , the long-term forecasts will follow a quadratic trend.

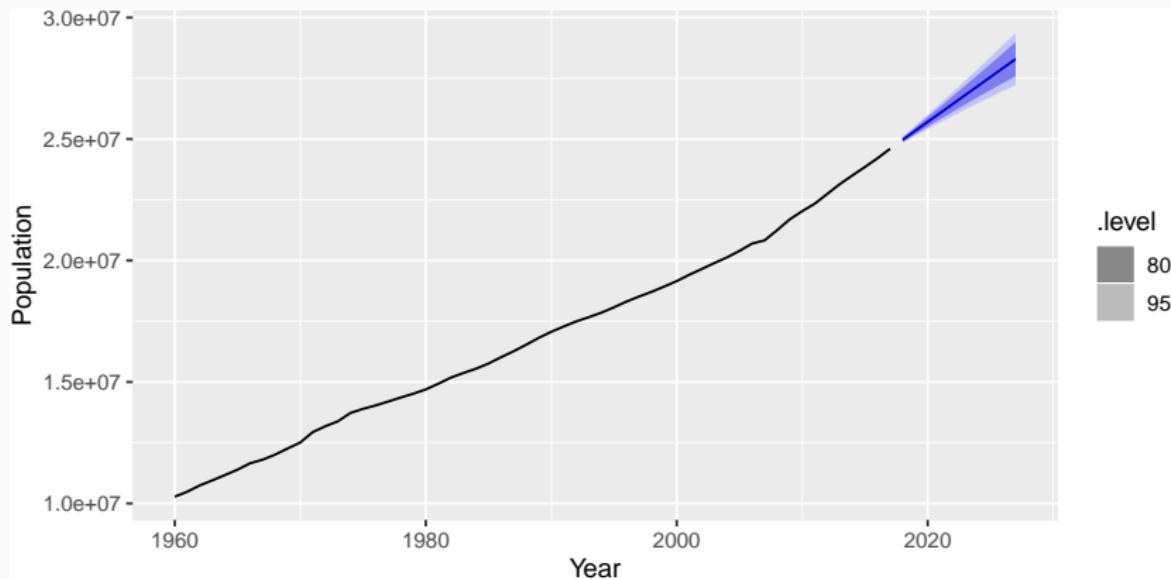
# Understanding ARIMA models

## Forecast variance and $d$

- The higher the value of  $d$ , the more rapidly the prediction intervals increase in size.
- For  $d = 0$ , the long-term forecast standard deviation will go to the standard deviation of the historical data.

# Example: National populations

```
fit %>% forecast(h=10) %>%  
  filter(Country=="Australia") %>%  
  autoplot(global_economy)
```



# How does ARIMA() work?

## Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences  $d$  via KPSS test.
- Select  $p, q$  and inclusion of  $c$  by minimising AICc.
- Use stepwise search to traverse model space.

# How does ARIMA() work?

## Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences  $d$  via KPSS test.
- Select  $p, q$  and inclusion of  $c$  by minimising AICc.
- Use stepwise search to traverse model space.

$$\text{AICc} = -2 \log(L) + 2(p+q+k+1) \left[ 1 + \frac{(p+q+k+2)}{T-p-q-k-2} \right].$$

where  $L$  is the maximised likelihood fitted to the differenced data,  $k = 1$  if  $c \neq 0$  and  $k = 0$  otherwise.

# How does ARIMA() work?

## Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences  $d$  via KPSS test.
- Select  $p, q$  and inclusion of  $c$  by minimising AICc.
- Use stepwise search to traverse model space.

$$\text{AICc} = -2 \log(L) + 2(p+q+k+1) \left[ 1 + \frac{(p+q+k+2)}{T-p-q-k-2} \right].$$

where  $L$  is the maximised likelihood fitted to the differenced data,  $k = 1$  if  $c \neq 0$  and  $k = 0$  otherwise.

Note: Can't compare AICc for different values of  $d$ .

# How does ARIMA() work?

**Step1:** Select current model (with smallest AICc) from:

ARIMA(2,  $d$ , 2)

ARIMA(0,  $d$ , 0)

ARIMA(1,  $d$ , 0)

ARIMA(0,  $d$ , 1)

# How does ARIMA() work?

**Step1:** Select current model (with smallest AICc) from:

ARIMA(2, d, 2)

ARIMA(0, d, 0)

ARIMA(1, d, 0)

ARIMA(0, d, 1)

**Step 2:** Consider variations of current model:

- vary one of  $p, q$ , from current model by  $\pm 1$ ;
- $p, q$  both vary from current model by  $\pm 1$ ;
- Include/exclude  $c$  from current model.

Model with lowest AICc becomes current model.

# How does ARIMA() work?

**Step1:** Select current model (with smallest AICc) from:

ARIMA(2, d, 2)

ARIMA(0, d, 0)

ARIMA(1, d, 0)

ARIMA(0, d, 1)

**Step 2:** Consider variations of current model:

- vary one of  $p, q$ , from current model by  $\pm 1$ ;
- $p, q$  both vary from current model by  $\pm 1$ ;
- Include/exclude  $c$  from current model.

Model with lowest AICc becomes current model.

**Repeat Step 2 until no lower AICc can be found.**

# Outline

- 1 What can we forecast?
- 2 Some case studies
- 3 The statistical forecasting perspective
- 4 Exponential smoothing
- 5 Lab Session 7
- 6 ARIMA models
- 7 Lab Session 8
- 8 Seasonal ARIMA models

## Lab Session 8

For the United States GDP data (from  
global\_economy):

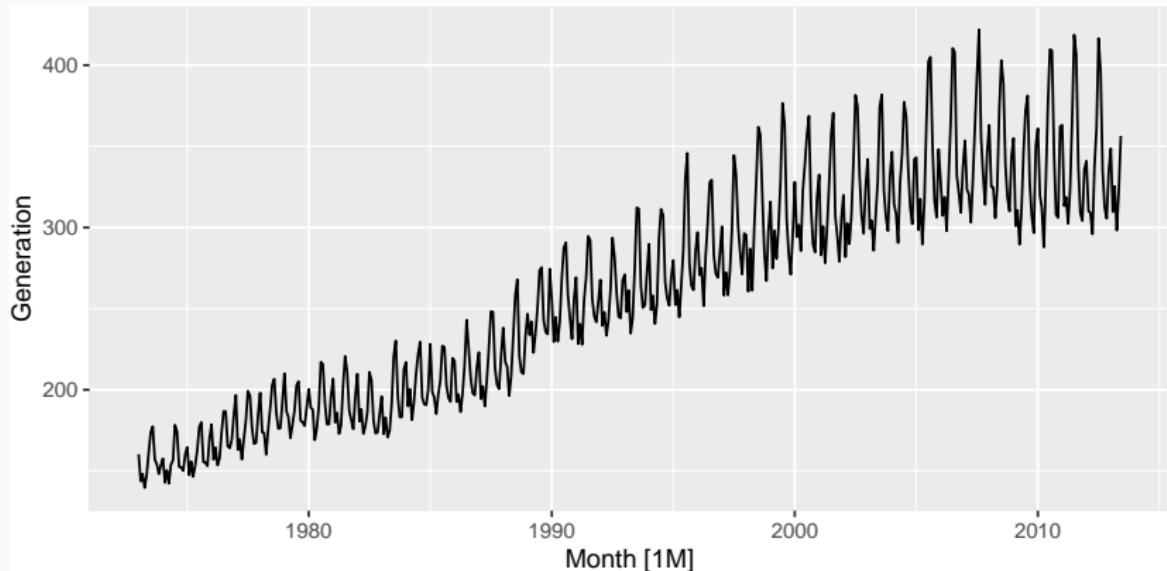
- Fit a suitable ARIMA model for the logged data.
- Produce forecasts of your fitted model. Do the forecasts look reasonable?

# Outline

- 1 What can we forecast?
- 2 Some case studies
- 3 The statistical forecasting perspective
- 4 Exponential smoothing
- 5 Lab Session 7
- 6 ARIMA models
- 7 Lab Session 8
- 8 Seasonal ARIMA models

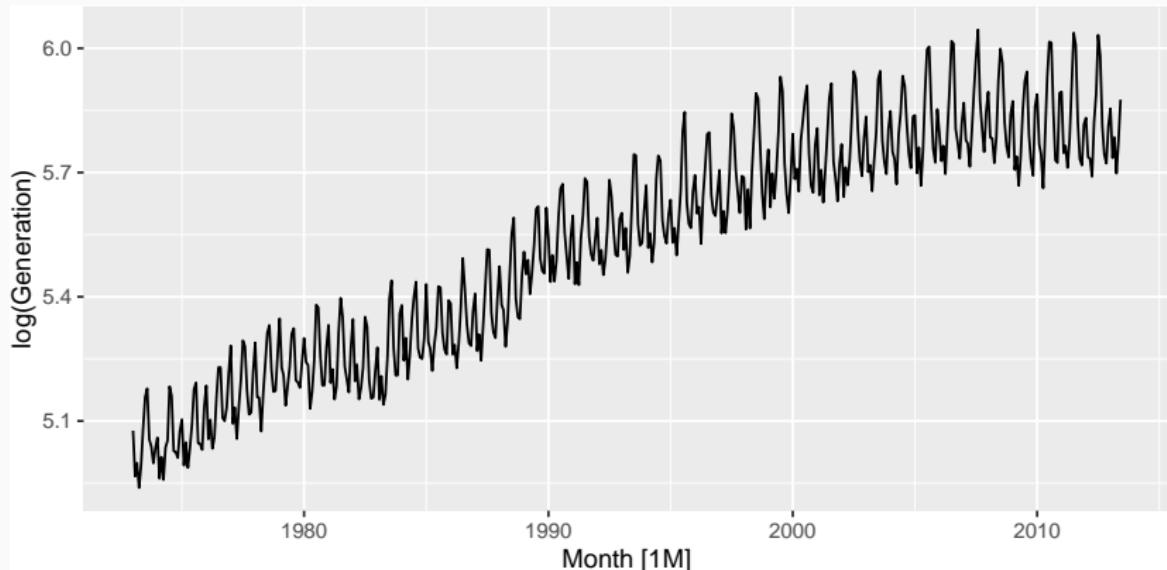
# Electricity production

```
usmelec %>% autoplot(  
  Generation  
)
```



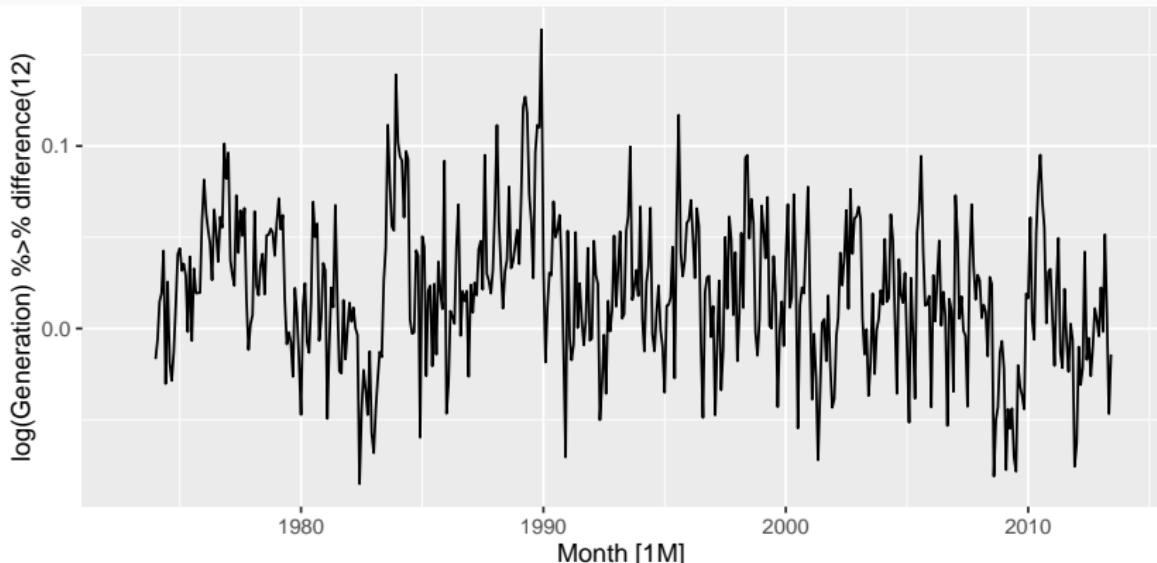
# Electricity production

```
usmelec %>% autoplot(  
  log(Generation)  
)
```



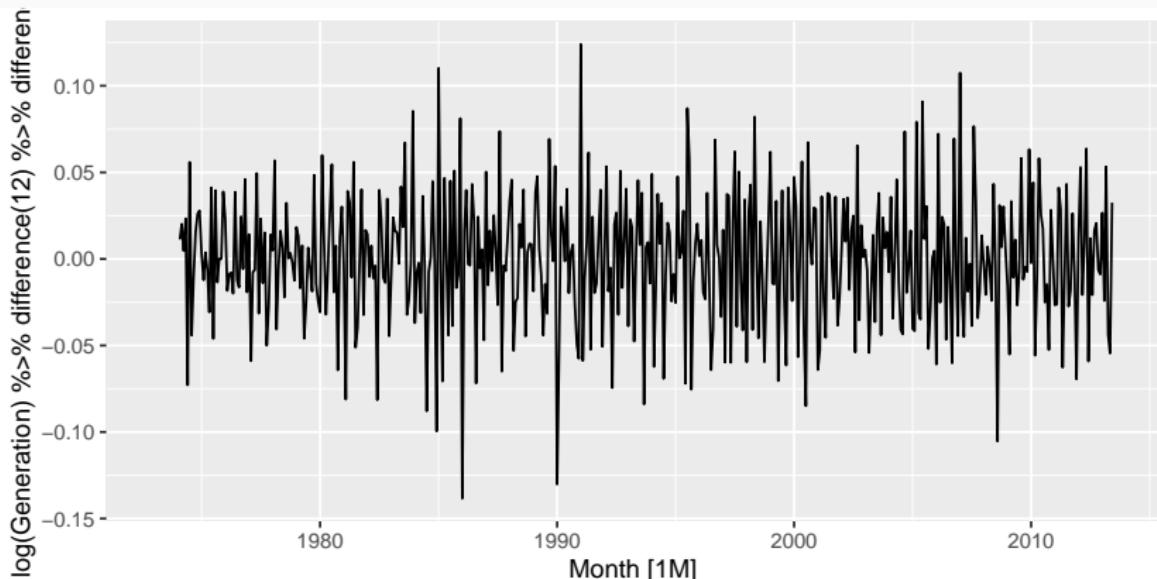
# Electricity production

```
usmelec %>% autoplot(  
  log(Generation) %>% difference(12)  
)
```



# Electricity production

```
usmelec %>% autoplot(  
  log(Generation) %>% difference(12) %>% difference()  
)
```



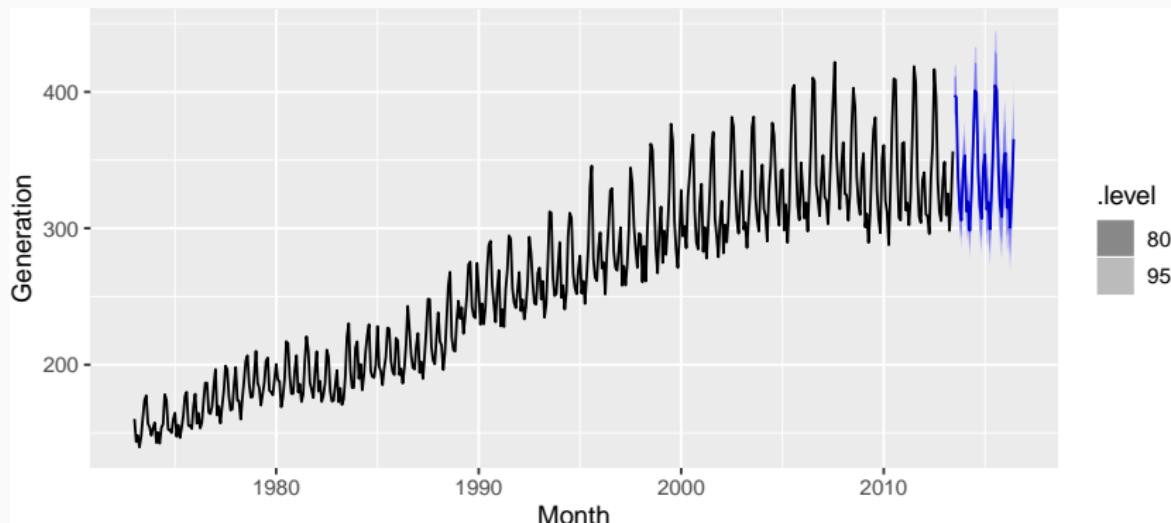
## Example: US electricity production

```
usmelec %>%  
  model(arima = ARIMA(log(Generation))) %>%  
  report()
```

```
## Series: Generation  
## Model: ARIMA(1,1,1)(2,1,1)[12]  
## Transformation: log(.x)  
##  
## Coefficients:  
##          ar1      ma1     sar1     sar2     sma1  
##          0.4116 -0.8483  0.0100 -0.1017 -0.8204  
## s.e.  0.0617  0.0348  0.0561  0.0529  0.0357  
##  
## sigma^2 estimated as 0.0006841:  log likelihood=1047  
## AIC=-2082    AICc=-2082    BIC=-2057
```

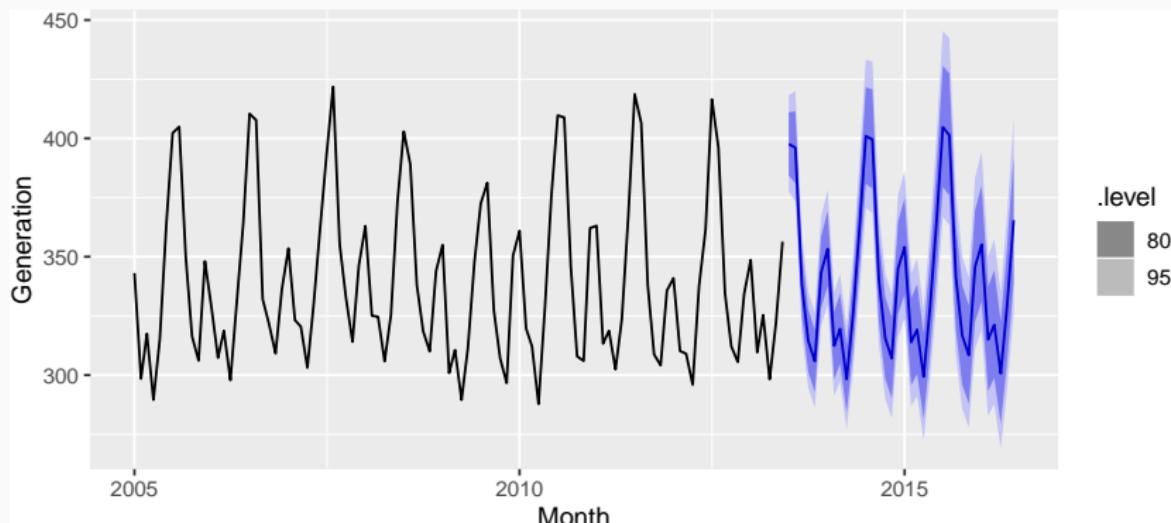
# Example: US electricity production

```
usmelec %>%  
  model(arima = ARIMA(log(Generation))) %>%  
  forecast(h="3 years") %>%  
  autoplot(usmelec)
```



# Example: US electricity production

```
usmelec %>%  
  model(arima = ARIMA(log(Generation))) %>%  
  forecast(h="3 years") %>%  
  autoplot(filter_index(usmelec, 2005 ~ .))
```



# Seasonal ARIMA models

ARIMA	$\underbrace{(p, d, q)}$	$\underbrace{(P, D, Q)_m}$
	↑	↑

Non-seasonal part      Seasonal part of  
of the model            of the model

- $m$  = number of observations per year.
- $d$  first differences,  $D$  seasonal differences
- $p$  AR lags,  $q$  MA lags
- $P$  seasonal AR lags,  $Q$  seasonal MA lags

Seasonal and non-seasonal terms combine  
multiplicatively

## Common ARIMA models

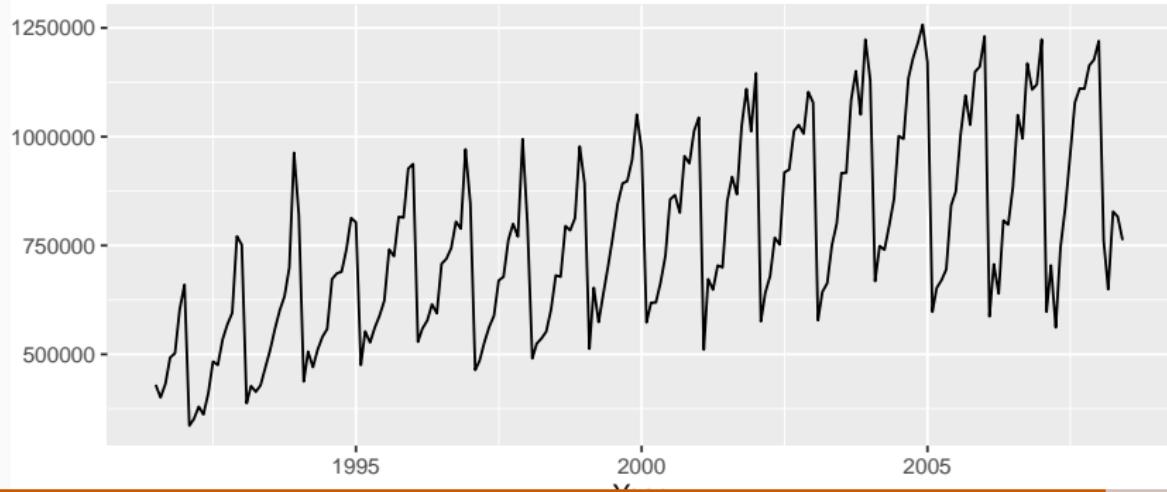
The US Census Bureau uses the following models most often:

ARIMA(0,1,1)(0,1,1) <sub>m</sub>	with log transformation
ARIMA(0,1,2)(0,1,1) <sub>m</sub>	with log transformation
ARIMA(2,1,0)(0,1,1) <sub>m</sub>	with log transformation
ARIMA(0,2,2)(0,1,1) <sub>m</sub>	with log transformation
ARIMA(2,1,2)(0,1,1) <sub>m</sub>	with no transformation

# Cortecosteroid drug sales

```
h02 <- PBS %>% filter(ATC2 == "H02") %>%  
  summarise(Cost = sum(Cost))  
h02 %>% autoplot(Cost) +  
  xlab("Year") + ylab("") +  
  ggtitle("Cortecosteroid drug scripts")
```

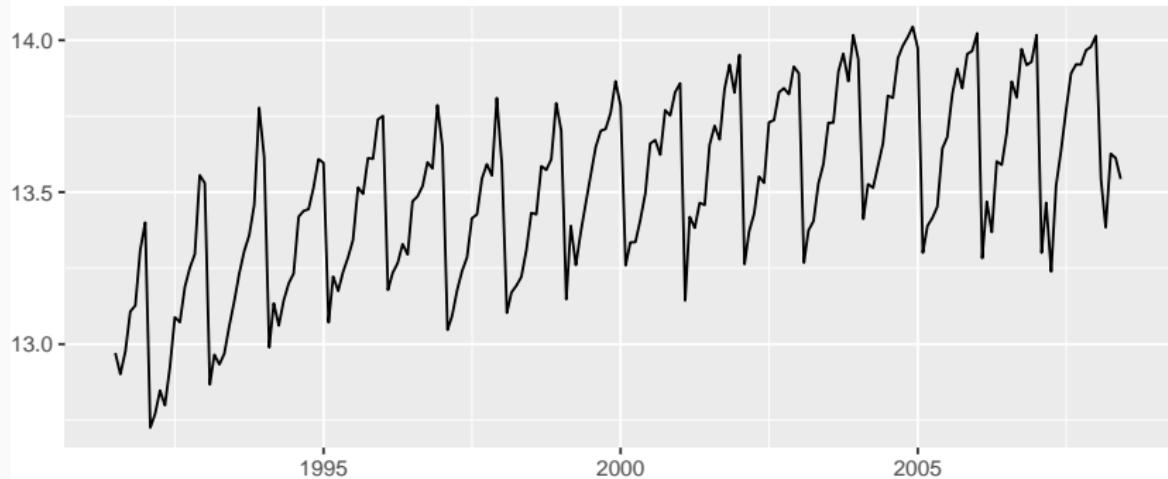
Cortecosteroid drug scripts



# Cortecosteroid drug sales

```
h02 <- PBS %>% filter(ATC2 == "H02") %>%  
  summarise(Cost = sum(Cost))  
h02 %>% autoplot(log(Cost)) +  
  xlab("Year") + ylab("") +  
  ggtitle("Log Cortecosteroid drug scripts")
```

Log Cortecosteroid drug scripts



# Cortecosteroid drug sales

```
fit <- h02 %>%
  model(auto = ARIMA(log(Cost)))
report(fit)

## Series: Cost
## Model: ARIMA(2,1,0)(0,1,1)[12]
## Transformation: log(.x)
##
## Coefficients:
##             ar1      ar2     sma1
##           -0.8491  -0.4207  -0.6401
## s.e.    0.0712   0.0714   0.0694
##
## sigma^2 estimated as 0.004399:  log likelihood=245
## AIC=-483  AICc=-483  BIC=-470
```

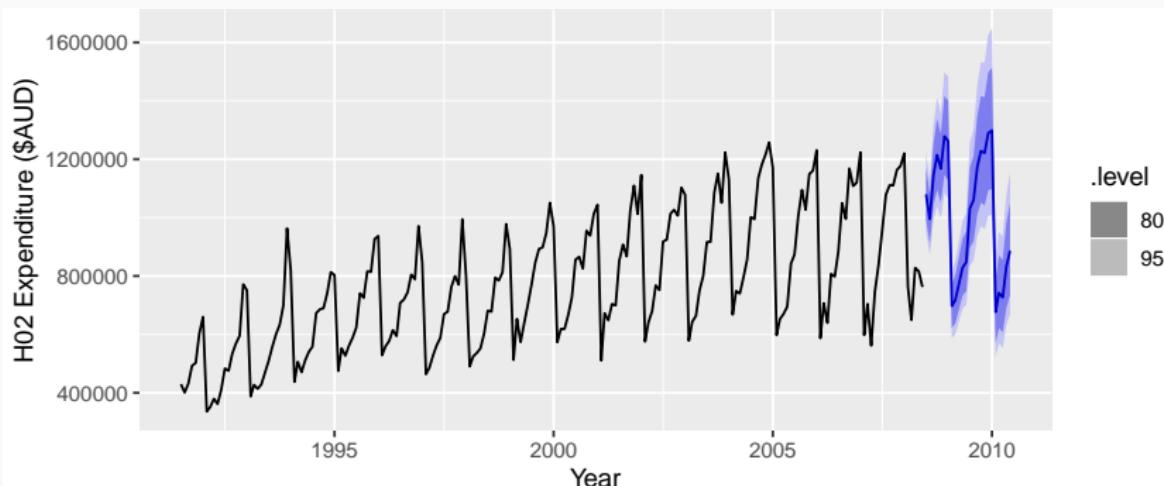
# Cortecosteroid drug sales

```
fit <- h02 %>%
  model(best = ARIMA(log(Cost), stepwise = FALSE,
                     approximation = FALSE,
                     order_constraint = p + q + P + Q <= 9))
report(fit)
```

```
## Series: Cost
## Model: ARIMA(4,1,1)(2,1,2)[12]
## Transformation: log(.x)
##
## Coefficients:
##             ar1     ar2     ar3     ar4      ma1     sar1     sar2
##             -0.0426  0.210   0.202  -0.227  -0.742   0.621  -0.383
## s.e.      0.2167  0.181   0.114   0.081   0.207   0.242   0.118
##             sma1     sma2
##             -1.202   0.496
## s.e.      0.249   0.214
##
## sigma^2 estimated as 0.004061:  log likelihood=254
## AIC=-489    AICc=-487    BIC=-456
```

# Cortecosteroid drug sales

```
fit %>% forecast %>% autoplot(h02) +  
  ylab("H02 Expenditure ($AUD)") + xlab("Year")
```



# Outline

1 What can we forecast?

2 Some case studies

3 The statistical forecasting perspective

4 Exponential smoothing

5 Lab Session 7

6 ARIMA models

7 Lab Session 8

8 Seasonal ARIMA models

# Lab Session 9

For the Australian tourism data (from tourism):

- Fit a suitable ARIMA model for all data.
- Produce forecasts of your fitted models.
- Check the forecasts for the “Snowy Mountains” and “Melbourne” regions. Do they look reasonable?

# Outline

1 What can we forecast?

2 Some case studies

3 The statistical forecasting perspective

4 Exponential smoothing

5 Lab Session 7

6 ARIMA models

7 Lab Session 8

8 Seasonal ARIMA models

# Training and test sets



- A model which fits the training data well will not necessarily forecast well.
- Forecast accuracy is based only on the test set.

## Forecast errors

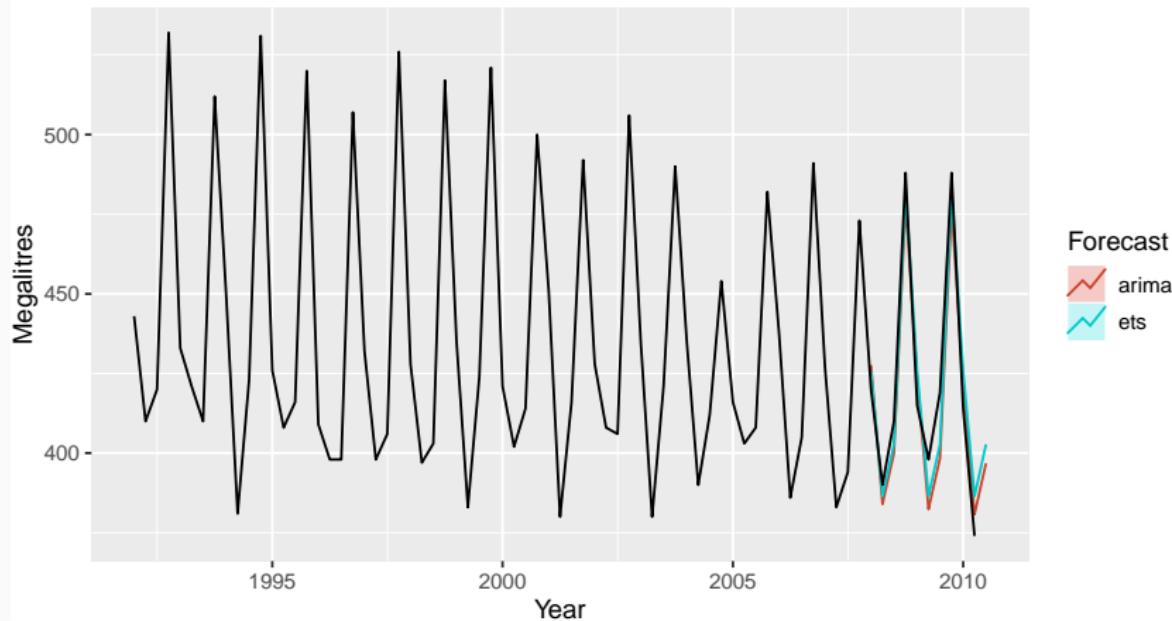
Forecast “error”: the difference between an observed value and its forecast.

$$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T},$$

where the training data is given by  $\{y_1, \dots, y_T\}$

# Measures of forecast accuracy

Forecasts for quarterly beer production



## Measures of forecast accuracy

$y_{T+h}$  =  $(T + h)$ th observation,  $h = 1, \dots, H$

$\hat{y}_{T+h|T}$  = its forecast based on data up to time  $T$ .

$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T}$

MAE =  $\text{mean}(|e_{T+h}|)$

MSE =  $\text{mean}(e_{T+h}^2)$

RMSE =  $\sqrt{\text{mean}(e_{T+h}^2)}$

MAPE =  $100\text{mean}(|e_{T+h}| / |y_{T+h}|)$

## Measures of forecast accuracy

$y_{T+h}$  =  $(T + h)$ th observation,  $h = 1, \dots, H$

$\hat{y}_{T+h|T}$  = its forecast based on data up to time  $T$ .

$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T}$

MAE =  $\text{mean}(|e_{T+h}|)$

MSE =  $\text{mean}(e_{T+h}^2)$

RMSE =  $\sqrt{\text{mean}(e_{T+h}^2)}$

MAPE =  $100\text{mean}(|e_{T+h}| / |y_{T+h}|)$

- MAE, MSE, RMSE are all scale dependent.
- MAPE is scale independent but is only sensible if  $y_t \gg 0$  for all  $t$ , and  $y$  has a natural zero.

# Measures of forecast accuracy

## Mean Absolute Scaled Error

$$\text{MASE} = \text{mean}(|e_{T+h}|)/Q$$

where Q is a stable measure of the scale of the time series  $\{y_t\}$ .

Proposed by Hyndman and Koehler (IJF, 2006).

For non-seasonal time series,

$$Q = (T - 1)^{-1} \sum_{t=2}^T |y_t - y_{t-1}|$$

works well. Then MASE is equivalent to MAE relative to a naïve method.

# Measures of forecast accuracy

## Mean Absolute Scaled Error

$$\text{MASE} = \text{mean}(|e_{T+h}|)/Q$$

where Q is a stable measure of the scale of the time series  $\{y_t\}$ .

Proposed by Hyndman and Koehler (IJF, 2006).

For seasonal time series,

$$Q = (T - m)^{-1} \sum_{t=m+1}^T |y_t - y_{t-m}|$$

works well. Then MASE is equivalent to MAE relative to a seasonal naïve method.

# Measures of forecast accuracy

```
recent_production <- aus_production %>%  
  filter(year(Quarter) >= 1992)  
train <- recent_production %>% filter(year(Quarter) <= 2007)  
beer_fit <- train %>%  
  model(  
    ets = ETS(Beer),  
    arima = ARIMA(Beer)  
  )  
beer_fc <- forecast(beer_fit, h="4 years")  
accuracy(beer_fc, aus_production)
```

```
## # A tibble: 2 x 9  
##   .model .type     ME   RMSE    MAE    MPE    MAPE    MASE    ACF1  
##   <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1 arima   Test   4.18  11.2  10.4  0.940  2.47  0.657  0.145  
## 2 ets     Test   0.854  9.80  8.99  0.151  2.18  0.568  0.207
```